Audit the chemistry you will screen: count-level guarantees for a screening round

Petrus H. Zwart*

Abstract

Conformal prediction returns prediction sets whose coverage is guaranteed, but in a screening campaign that guarantee is a request rather than a finding: the level is chosen and the threshold is built to deliver it, so every setting is equally guaranteed and coverage ranks none of them. It cannot say which setting to screen under, how many compounds will be assayed, or how large a round must be. We keep the two class-wise requests as coordinates instead and state what each one does as counts of the outcomes a campaign is planned on: confident hits, false leads, missed in-class compounds, compounds assayed. Expectations of those counts cannot be planned on or constrained with a stated probability, so we bound them in finite samples from a labelled audit set drawn from the library to be screened, and invert the bound to read off the round size an audited count implies. The audit is a second step of the same form, exempting the calibration set as the first exempts the training set, so its counts survive a calibration that does not match the screen; what it requires instead is the chemistry the batch will contain. In a DNA-encoded library screened for soluble epoxide hydrolase, an audit matching only the batch's proportion of new chemistry covered 0.63 ± 0.10 of the consequence lines, no better than the 0.63 ± 0.13 of ignoring the shift altogether and no nearer the request; the two differ by 0.000 ± 0.070 on a paired comparison. An audit drawn from the batch's own building blocks covered 0.97 ± 0.04, but that arrangement splits one pool at random and so serves as a positive control rather than as evidence that a pilot repairs a broken campaign. Withholding blocks, the split used to guard against memorisation, cost 0.07 in ranking and two thirds of the coverage, so the ranking a screen reports does not announce the loss. Whether the match holds is testable beforehand from structures alone: the two preparations that fail read 0.66 against a permutation null centred on 0.50, and the one that holds reads 0.50. Dividing a pooled electrolyte compilation by the paper each measurement came from, rather than at random, drops coverage from 0.97 ± 0.04 to 0.70 ± 0.12, so provenance is part of what must be exchangeable. A batch the model chose is the hardest case and the one a self-driving laboratory commits: an audit reserved from the library could not even be formed in nine of twelve rounds, because nothing in-class was left to sample, while a random slice of the committed batch covered every line in all nine rounds where it could be. One round of 1500 candidates from a pool of 70 000 also took the in-class rate of the remainder from 0.0071 to 0.0013, so an audit prices the round it was drawn from and not the one after. Buying a pilot of the screened source repairs the provenance case: audited off-source, coverage is 0.55 ± 0.07; audited from a pilot of that source divided by concentration, by solvent system, or as the paper reported it — divisions that are not random and so can fail — it is 0.86 to 0.92. Quoted spreads are two standard errors over repeats, the repeat being the unit.

1 Introduction

A screening team makes two splits, and believes the wrong one is the careful one. It separates training compounds from the rest by scaffold or building block, to check that a model has not simply memorised its training chemistry; that is the split a paper reports, and a ranking that survives it is read as a licence to deploy. It then divides what remains at random to calibrate a rule. The first split looks rigorous and the second looks like bookkeeping.

Both readings are wrong, and in opposite directions. The structural split does guard the model, and it silently voids everything quantitative about the next batch: in a DNA-encoded library, withholding building blocks cost two thirds of the coverage of the predicted counts while costing 0.07 in ranking, a drop no project would act on. The random division, meanwhile, is what keeps those counts valid — but only for the population it sampled, and that is seldom the population a round commits. When a planner fills a plate from the top of a ranking, an audit reserved from the library does not price it; in nine of twelve rounds of the campaign below such an audit could not even be formed, because nothing in-class remained in the leftover pool to measure. The split a team trusts protects the wrong thing, and the split it treats as housekeeping protects the right thing for the wrong population.

This paper is about that second split: what a labelled sample licenses a team to claim about a round, and what each way of drawing it wrong costs.

The claims in question are counts, and they have to be. A round is constrained in integers — the wells on a plate, the reagent budget of a run, the confirmed hits needed to justify the next round — and those constraints cannot be expressed as rates. A rate slightly off is invisible; a plate that overruns its capacity is not. Conformal prediction, which wraps any trained model and returns prediction sets with guaranteed coverage under exchangeability,[1,2] does not supply them: coverage is a level the team chooses and the threshold is built to deliver, so every setting is equally guaranteed and coverage ranks none of them. Risk-controlling methods bound a chosen loss, but what they control is its expectation or its level[3–5] rather than the integer a plate returns, and the post hoc literature that does bound counts describes a set already in hand (Section 1). What is missing is a count that is *predictive*: a statement about a batch not yet made. That is what a labelled audit set supplies, and why the rest of this paper is

about how the audit is drawn.

The screens we consider decide which candidates go forward to a toxicity assay, a solubility or conductivity measurement, or synthesis. Because characterization runs in parallel, the capacity of a round is fixed before it starts. Unlike turn-by-turn optimization, a batch campaign often aims at covering a chemical or physical design space, mapping it or narrowing in on a set of optima (Section 4). The decision, whether made by a team or by a planner, is the same, and so is the gap we address. In a binary screen each compound receives one of four prediction sets: {out-class}, {in-class}, both labels, or neither. A team acts on these sets. It advances or discards compounds that receive a single label and sends the others for an assay.

To choose among these actions, a team needs to know what they will do: how many compounds will receive a single label, how many of those calls will be right, and how many compounds will be sent for testing. These consequences, counted in compounds, are the terms in which screening decisions are made and budgets and capacities are planned. Coverage, the native coordinate of conformal prediction, is a statement about the prediction sets; it informs a decision once it is linked to these consequences. That link has to be measured. A set containing both labels is always correct but sends the compound to an assay, so two classifiers calibrated to the same coverage can sort a batch very differently.

Predictive uncertainty is ordinarily reported one prediction at a time, and reviews of uncertainty quantification for drug design and property prediction follow that convention.[6] Here the unit of prediction is the round rather than the compound, which changes what the uncertainty is made of. Two terms contribute to it: how well the audit pins down the rates, and how much the realized count varies around those rates in a batch of the size committed. The second does not shrink however large the audit grows, and at the batch sizes a plate imposes it is the larger of the two — which is why a guarantee stated as a rate does not answer the question, and why the construction below is stated in counts.

What we do and why. We give coverage a specific role. The coverage level is a setting the team chooses, separately for the in-class and out-class compounds, and changing it changes which compounds are advanced, discarded or assayed. We therefore treat the coverage request as the control: sweeping it moves the classifier through the space of possible decisions. A labelled *audit set*, a random sample of the library to be screened that was used neither for training nor for calibration, shows the consequences of each setting. From it we predict every screening outcome of a future batch, stated as counts of compounds, with prediction intervals that are valid in finite samples under exchangeability of the audit with the batch. That is a different requirement from the one coverage assumes, and Section 2.5 separates them: either can hold without the other.

These intervals are the deliverable, and obtaining them is what the method is for. Before a batch is screened, the audit states how many confident hits, false leads, missed in-class compounds and measurements the batch will contain, each as a count, equivalently a rate per 1000, with an interval that holds with a specified probability. The intervals are realistic in both directions: they hold at the stated level, where the usual shortcuts are too narrow, and they are no wider than the size of the audit requires. A team can therefore plan on them: book the assays, budget the measurements and expect the hits.

Coverage is thus the coordinate for navigating decisions, and the audited consequences are what a team chooses by. Choosing uses an objective, such as the expected cost or yield of the screen. It usually also uses guarantees: a plan the team can rely on, in which, with a stated probability, no batch needs more assays than the laboratory or the robot can run and every batch delivers enough hits to proceed. Such guarantees are bounds on counts that hold with a stated probability. The coverage requests reduce the navigation of these many consequences to two coordinates, in which a team can explore scenarios, each a set of such guarantees, and see which settings satisfy them.

Section 2 sets out the outcome table, the threshold correction, the navigation by coverage and the audit intervals that are the method's output; Section 3 validates the intervals in simulation and on real data and applies them to solubility, toxicity and electrolyte screens; Section 4 gives a workflow and the limitations.

Contributions. The statistical tools are classical; what is new is the separation they make visible and what violating it costs.

- *A labelled audit set prices a round, in counts, before it is run.* Sweeping the two class-wise coverage requests traces the decisions available to a team, and at each setting an audit gives finite-sample intervals for every cell of the outcome table — confident hits, false leads, missed in-class compounds, compounds assayed — valid after the setting is chosen and invertible for the size of the round.[7,8]
- *Two guarantees, two requirements, and the asymmetry runs one way.* From the audit's point of view the calibration set is training data, so coverage and predicted counts fail independently (Table 2); a rule calibrated on the wrong population, prevalence or label is a worse rule but a valid audit, while an audit drawn from chemistry the batch does not share voids the counts however the rule was calibrated.
- *What violating it costs, measured four ways on real data.* Withheld building blocks, division by provenance, an audit matching a batch's composition but not its chemistry, and a batch the model chose: the counts break in all four, and no statistic a screen reports orders the failures.
- *What satisfying it takes.* A pilot of the screened population repairs it under divisions that are not random, and a random slice of a committed batch repairs it exactly, whatever rule selected the batch. One round of selection then depletes the library, so the slice is re-drawn each round rather than amortized.
- *A check that costs no measurements.* Whether the audit matches the batch is testable beforehand from structures alone, as a classifier two-sample test against a permutation null[9] — one-sided, so it refuses an arrangement rather than certifying one, and blind by construction to a shift the representation cannot see.

Related work. Conformal prediction is established in cheminformatics, for toxicity, bioactivity, imbalanced data, applicability domain and drift.[10–19] The closest work estimates what a screening round will return. Buendia *et al.* predict the number of hits an iteration of an iterative high-throughput screen will retrieve, by summing per-compound calibrated probabilities from Venn–ABERS predictors over the selected compounds, and show across six targets that the sum tracks the realized count;[20,21] Svensson *et al.* link conformal predictions to the gain of a screen.[22] Estimating campaign-level counts is therefore established, and we build on it. A sum of calibrated probabilities is an expectation, and we add a probability statement about the realized count, which is what a plan that has to hold requires. Buendia *et al.* also identify the obstacle to such a statement: selecting the highest-ranked compounds leaves the calibration and screened sets progressively less exchangeable, retaining validity would require measuring a fresh random subset of the screened library at every iteration, and the cost of those measurements would make the approach impractical. They therefore keep the initial calibration set, and report that the realized hit count falls outside their expected limits in later iterations. That randomly chosen, measured subset is the audit set used here, and its cost can be set against the compounds it saves (Section 3.3). These studies report coverage and efficiency, or a gain, on test data; we add prediction intervals for what a future batch will contain. In statistics, the coverage of a conformal predictor given its calibration set is known exactly,[23–25] and related to tolerance regions and PAC prediction sets.[26,27] Risk-controlling methods are the closest prior art and the comparison is worth making precisely, because they answer a different question rather than a weaker version of the same one. Conformal risk control bounds an expectation, $\mathbb{E}[L_{n+1}(\hat{\lambda})] \leq \alpha$, marginally over the calibration set and the test point together, for a single future point, under a loss required to be monotone in a scalar parameter.[4] The loss itself is general — it may be any bounded function, including an indicator that a plate has overrun its capacity, so the family is not confined to per-candidate error rates — but what is controlled is its expectation or its level, not the integer the plate returns. Risk-controlling prediction sets make that statement high-probability and conditional on the calibration set, $\Pr(R(\hat{\lambda}) \leq \alpha) \geq 1-\delta$;[3] learn-then-test extends it to several parameters and several risk constraints.[5] Song *et al.* come closest to the construction here, bounding the false discovery proportion of a conformal selection with a bound that holds simultaneously over all thresholds, so that a threshold may be chosen after the data are seen;[28] their envelope is a sharper instrument than the Bonferroni correction we use for the same purpose, and the quantity it bounds is a proportion among the selections of the batch in hand rather than a count in a batch not yet measured. Bounding a realized count in a finite selected set is older than any of this and is not ours: the post hoc, true-discovery literature has given simultaneous high-probability bounds on the number of true discoveries in an arbitrary selected set since 2011.[29–31] What distinguishes the construction here is not that it counts but that it is *predictive*: those bounds describe the set in hand, from p-values computed on it, whereas the intervals below describe a batch not yet made or measured, from a sample of the library it will be drawn from. They are also two-sided, so a capacity and a yield constraint can be imposed at once; they cover every cell of the outcome table, including the missed in-class compounds, which a bound computed without labels cannot reach because a miss is defined by a label it never sees; and they invert for the batch size. The price is labels: an audit is measured, and the post hoc bounds are not.

A campaign is constrained in counts, not rates, and the two differ by more than a rounding. A guarantee that the long-run failure rate is at most 0.1 does not say that a plate of 96 will hold at most 13 failures with probability 0.95, because the realized count varies around the rate by an amount the rate bound does not describe. That variability is not a correction at the scale these platforms run: at an audit of 400 and a batch of 50, 89% of the variance of the interval reported here is contributed by the draw of the batch rather than by uncertainty in the rate (Section 3.3). A bound on the rate alone therefore leaves the larger part of that variability undescribed, and leaves most of it undescribed exactly where autonomous platforms operate, in small batches run often. A rate bound can of course be composed with a binomial tail to yield a count bound, and the composition is valid. It pays twice, once for the rate and once for the tail, where conditioning on the audit-and-batch total pays once, and the price is measurable. Giving the rate bound the most generous instrument available — an exact Clopper–Pearson limit rather than the concentration bounds risk-controlling sets actually use — and matching both sides at a one-sided level of 0.05, the composed bound is 11 to 60% wider than the exact inversion over the configurations this paper reports, and the gap is largest where it matters most: pricing one plate of 96 from an audit of 384 at a 2% rate, the exact bound allows 5 confident hits and the composition 8, so the headroom a planner must reserve doubles. The construction here bounds the realized integer count in a batch of stated size, two-sided so that capacity and yield can both be constrained, for every cell of the outcome table rather than one monotone summary, without any monotonicity requirement, and invertible for the batch size itself. Conformal selection[32] selects candidates with false-discovery control, and set-valued classifiers with bounded class-wise error[33] minimize ambiguity. Decision-theoretic work links prediction sets to actions. Conformal decision theory calibrates a decision parameter to a target risk,[34] and Kiyani *et al.* showed that prediction sets are a sufficient statistic for risk-averse decision makers with a known utility, for whom acting on the worst case within the set is optimal.[35] The question for which costs a classifier, or an abstaining classifier, is optimal goes back to the reject option[36] and to ROC analysis.[37,38] Our aim is complementary. We assume no utility, report *all* screening outcomes of each coverage setting with valid intervals and bounds, and let a team choose afterwards under its own objective and constraints.

2 Methods

2.1 Setting and assumptions

We call the class with the property the screen looks for *in-class* (label 1) and the other *out-class* (label 0): compounds active at an endpoint in Tox21, soluble compounds in AqSolDB, and electrolytes that reach the conductivity target in CALiSol-23. A trained binary classifier is held fixed. It gives each compound a predicted probability of being in-class, $u = \hat{p}(\text{in-class} \mid x)$. A labelled *calibration set* sets the

thresholds. A labelled *audit set*, which in machine-learning terms is a held-out test set, is used only to measure what the thresholds do; we call it an audit set throughout, because what it is for here is pricing a round rather than scoring a model. A *future batch* of m compounds will be screened. We assume throughout that the audit and future compounds are exchangeable with one another given the trained model and the calibration set, *i.e.* that their joint distribution is unchanged when they are reordered. The requirement is conditional on the calibration set, which need not itself be exchangeable with either; a random division of a fixed library satisfies it even though the three parts are then dependent through the library's finite size. Where the coverage guarantee is also wanted, the calibration set must in addition be exchangeable with the batch; the two requirements are separated in Section 2.5 and they fail independently. What must be exchangeable is each candidate together with the label recorded for it, not the candidate alone, so the assumption can fail in two ways that look nothing alike: the compounds may be drawn from different populations, or the compounds may be interchangeable while the label recorded on one set is not the label recorded on another. The first is a question of which compounds a round contains and the second of what is measured on them, and both appear below. Exchangeability holds when all three sets are random samples of the library to be screened and carry the same label. The training set is deliberately absent from that list. The classifier is held fixed and enters only as a score, so how its training compounds were chosen bears on how good that score is and not on whether the intervals are valid. The two splits a campaign makes are therefore independent and answer different questions. Separating the training compounds from the rest by scaffold, series or building block guards against a model that has memorised its training chemistry, and lowers an optimistic estimate of how well it ranks compounds unlike the ones it was fitted on. Dividing what remains into calibration, audit and batch at random is what keeps the intervals valid. Both can be done at once, and should be: the arrangement that measures a model on unseen chemistry *and* prices the batch correctly is a split at the training boundary followed by random assignment within the deployment population. The unit of sampling is the candidate to be screened: a compound, or for electrolytes a formulation, not an individual measurement. It fails when the batch comes from other chemical series or scaffolds (Section 3.12), and again when the thresholds are set against a cheaper label than the one the campaign's claim concerns (Section 3.11). It fails too when the batch is chosen using the model, as in active learning, unless that selection is part of the design. Section 3.9 tests that case and the two repairs available for it. The probabilities need not be calibrated in the usual sense; only their ordering matters.

2.2 Prediction sets and the outcome table

We use class-conditional (Mondrian) split conformal prediction,[1] which controls coverage separately for the in-class and out-class compounds. The label *out-class* is included in a compound's prediction set when $u \le \tau_0$, where τ_0 is set from the out-class calibration compounds. The label *in-class* is included when $u \ge 1 - \tau_1$, where τ_1 is set from the in-class calibration compounds (Fig. 1; ESI Section S1). Each compound then falls into one cell of a 2×4 table of true label by

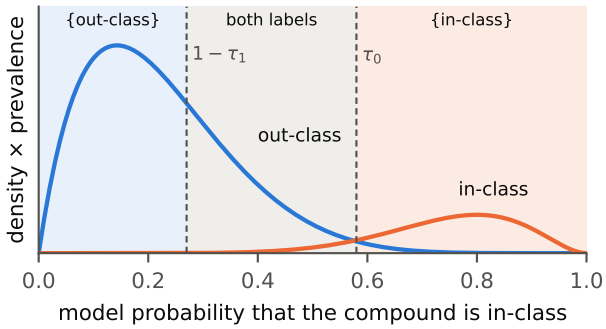


Figure 1: Class-conditional thresholds divide compounds into prediction sets. Curves show the distribution of the predicted probability of being in-class for out-class (blue) and in-class (orange) compounds, weighted by prevalence. Compounds left of $1 - \tau_1$ receive {out-class}, those right of τ_0 receive {in-class}, and those in between receive both labels. If the thresholds cross, the middle band receives the empty set instead. The coverage request for each class moves one threshold.

Table 1: Outcomes of a binary conformal screen. Each compound falls into one cell; every screening quantity is a sum of cells.

True label	{out-class}	{in-class}	both	none
out-class	correct rejection	false lead	deferred	deferred
in-class	missed in-class	confident hit	deferred	deferred

prediction set (Table 1). Every screening quantity is a sum of cells:

- missed in-class compounds are in-class compounds labelled {out-class};
- false leads are out-class compounds labelled {in-class}, *i.e.* compounds followed up that turn out out-class;
- deferred compounds receive both labels or none;
- coverage is the sum of the cells in which the true label is in the set;
- the hit rate of the confident calls is a ratio of two cells.

The cells in each row add up to the number of compounds of that class. Moving a threshold therefore moves compounds between cells of the same row, so the outcomes trade off against each other. With class probabilities, a compound can receive both labels or none, but not both kinds for the same thresholds (ESI Section S6). We report the two kinds separately, because an empty set signals a compound unlike the calibration data rather than an ambiguous one.[10]

2.3 Coverage requests and thresholds

The standard threshold guarantees coverage of at least $1 - \alpha$ only on average over calibration sets. With tens of in-class calibration compounds, the coverage of the one model actually deployed can fall well below $1 - \alpha$. The coverage observed over a batch of m_y future compounds of class y follows a Beta–Binomial distribution that depends only on the number of calibration compounds of that class and on the chosen order statistic. This result requires only exchangeability. It is the classical distribution-free rank law behind Wilks–Wald

tolerance limits, read in the conformal setting;[23,25,26] the correction rule built from it is a two-parameter tolerance limit transposed to a finite batch. We choose the least conservative threshold for which batch coverage is at least $1-\alpha_y$ with probability at least $1-\delta_y$, separately for each class. A coverage request therefore has four numbers, $(\alpha_0, \delta_0, \alpha_1, \delta_1)$. The guarantee is stated for a planned number m_y of future compounds of each class, which we set from the batch size and the class proportions of the calibration set. That is a plug-in for a quantity the campaign does not know, and at low prevalence it is estimated from few compounds, so the tolerance δ is controlled for the planned count rather than the realized one and the deployed guarantee is approximate in it; the correction is conservative in the count, and the real-data checks below came out conservative, but we have not bounded the error. In an automated laboratory the batch is a plate or one run of the platform, so m is set by its capacity. The batch-free version,[23] which follows a Beta law, controls coverage only for an infinitely large batch (Section 3.1). As baselines we use the standard threshold and a one-sided correction based on the Dvoretzky–Kiefer–Wolfowitz inequality.[39]

Tied scores, for example from identical fingerprints, keep the coverage guarantee conservative, but they can break the count intervals that take a threshold from the same compounds they then count, because tied compounds cross it together; the audit intervals are unaffected, since the thresholds are fixed before the audit is counted. We therefore break ties at random: every compound receives an independent uniform number, and compounds are ordered by score and then by that number. This needs no knowledge of future scores and makes all the laws exact (ESI Section S2).

2.4 Navigating decisions with the coverage request

Coverage as a coordinate. With class probabilities, every rule that acts on single labels and defers the rest is a two-threshold rule on u: label out-class below one threshold, in-class above another, defer in between. The two coverage requests set the two thresholds, so a sweep of (α_0, α_1) over a grid, at fixed δ, traces the available two-threshold rules, up to the resolution of the calibration scores and of the grid. We keep one representative for each distinct pair of thresholds. Each setting maps to a vector of consequences, the counts of the outcome table and their bounds; the sweep traces a two-dimensional surface in this many-dimensional space, and its Pareto front over the outcomes that matter is the map on which a team explores its options.

The decision geometry of a setting. Without stated costs, a setting together with the usual rule (act on single labels, defer the rest) induces a decision geometry: two thresholds on the model score. Given only the prediction sets, this rule is the risk-averse (max–min) optimal use of the sets for any costs in which a deferral is cheaper than either error (ESI Section S7).[35] Its geometry matches Chow's cost-optimal reject rule[36] when the probability of being in-class increases with u, so each setting is expected-cost optimal for some region of cost ratios, possibly empty. Which costs a setting serves depends on the fractions of in-class compounds within each prediction set, which labelled data reveal; for a known utility, sets built for that utility can serve it better.[35]

Table 2: The two guarantees are separate, and the asymmetry has a direction. Each cell names the experiments in this paper that occupy it. A rule calibrated on the wrong population, at the wrong prevalence or against a cheaper label is a worse rule but still a valid audit; an audit drawn from chemistry the batch does not share voids the counts however the rule was calibrated. The remaining cell needs a calibration set that matches the batch while the audit does not, which is an arrangement no design here produces and no team would choose deliberately, though it is what reusing a stale audit against fresh calibration would give.

	counts hold	counts fail
coverage holds	random assignment (53 of 54 lines)	not tested (see caption)
coverage fails	a new salt (29% coverage, audit 95.5–100%); a structured proxy (0.721 against 0.90, audit 95.2%)	withheld building blocks (AUROC 0.892, coverage 35%); division by provenance (0.97 to 0.70)

The coverage request therefore locates a decision, and its audited consequences evaluate it.

Objectives and chance constraints. A team chooses a setting by an objective, such as the expected number of false leads, the expected cost or the expected yield, subject to constraints on what the campaign may consume or must deliver: at most U assays, at most U' missed in-class compounds, at least L confirmed in-class compounds in a batch of m. Each constraint must hold with probability at least $1-\varepsilon$, so it is checked against a one-sided bound from the audit set, not against the expected count. This is a chance-constrained program[40] with two decision variables, solved by enumerating the grid, in which the constraint probabilities are replaced by distribution-free bounds from the audit. A *scenario* is such a set of constraints. Its feasible settings form a region of the coverage plane, and the objective selects a setting within it. When the objective and every constraint are monotone in the outcomes (fewer measurements, false leads or misses, and more confirmed compounds, are never worse), an optimal setting always lies on the Pareto front of the audited outcomes, taken over the quantities the objective and constraints use: a dominated setting is never better than the setting that dominates it and meets no constraint that setting fails. For every such objective, the table of Pareto settings with their expected counts and bounds is therefore a sufficient summary, and no cost ratio is needed to decide. Only an objective that is not monotone, such as a target workload that should be neither exceeded nor undershot, needs the full map. Changing the scenario, for example tightening the assay budget, moves or shrinks the region and shows what the change costs. An empty region is itself a result: the requirements cannot be guaranteed with the labelled data in hand, and either a constraint must be relaxed or more compounds labelled.

2.5 Prediction intervals for the consequences

What the method returns. For a chosen setting and a planned batch of m candidates, the output is one line per outcome of Table 1: the expected count, and an interval that contains the realized count with a stated probability. Counts for a batch of 1000 are rates per 1000, so the same statement can be read either way. Tables 4 and 5 show the form. Every-

thing else in this paper is built from these intervals: the sweep compares settings by them, the scenarios constrain them, and the audit-size rule plans their width.

From the audit set. Suppose an outcome (one cell, or a sum of cells) occurs K times among n audit compounds. Given the calibration set, the audit compounds and the m future compounds are exchangeable. Given the total count, the number in the audit set is therefore hypergeometric, whatever the data distribution. Inverting this exact hypergeometric test gives conservative limits $[L,U]$ for the future count K_m that hold with probability at least $1-\gamma$; they are quantiles of two Beta–Binomial distributions[7] (ESI Section S3). We use $\gamma = 0.05$ throughout. The construction holds for any outcome, any score and any fixed decision rule; it does not depend on conformal prediction. In the language of quality control the audit set is an acceptance sample, and the limits are the classical tolerance limits for the number of defectives in a future lot; the same logic, of bounding a finite population's count from a measured sample and inverting for how large that sample must be, is what makes a risk-limiting election audit work.[41,42 26,43]

Two steps, two requirements. It is worth naming what that conditioning does, because it accounts for the results below. Split conformal fits a score on training data, freezes it, and asks that calibration and future candidates be exchangeable; the training set is exempt precisely because it enters only through the frozen score. The audit repeats the move one level up. The rule it evaluates — the score together with the thresholds — is fixed by the training and calibration sets, and is then frozen, so the audit asks that audit and future candidates be exchangeable and exempts both of the sets that built the rule. From the audit's point of view the calibration set is training data. The argument is the same one twice, though not the same formula: a rank is uniform and a split of a count is hypergeometric for the one reason that under exchangeability every ordering, and so every subset of a given size, is equally likely.

Each step therefore carries its own guarantee and its own requirement, and a campaign can meet either without the other. Coverage is the first step's guarantee and needs the calibration set to be exchangeable with the batch. The predicted counts are the second step's and need the audit set to be, whatever the calibration was built from. So a rule calibrated on the wrong population, at the wrong prevalence or against a cheaper label is a worse rule but not an invalid audit, while an audit drawn from chemistry the batch does not share invalidates the counts however well the rule was calibrated; both appear below. The same reading explains why the calibration set cannot freely audit itself: for the second step it is training data, and a rule may not be priced on the data that set it (Section 3.12).

What the intervals cover. The intervals cover the realized count in the future batch; no fixed underlying rate is assumed. They are conservative because counts are discrete. Each interval is valid for one outcome and one setting fixed before the audit set is seen. Choosing a setting by looking at its audited outcomes and then quoting its intervals is optimistic. Intervals computed at level γ/J for each of the J requests in the grid hold simultaneously, and therefore remain valid for whichever setting is chosen. For a scenario with C bounded outcomes, each one-sided bound is computed at level $\varepsilon/(JC)$, so that every bound of every setting holds together with probability at least $1-\varepsilon$ and the chosen setting satisfies all constraints with at least that probability. Because the bounds hold jointly for all settings and all bounded outcomes, a team may pose any scenario over these outcomes, and change it, after looking at the map, without losing the guarantee. All intervals and bounds are conditional on the calibration set, that is, on the thresholds actually deployed.

What makes an interval useful. Validity alone is attained by the trivial interval $[0,m]$, which contains every possible count of an outcome in a batch of m and says nothing. An interval therefore earns its place by its width, and we report how many times narrower it is than that baseline, $m/(U-L)$, alongside coverage, following the principle of sharpness subject to calibration.[44] Width is set by the audit size and by how well the model ranks, and it can be planned before any labelling (Section 4). Whether the rule is better than no model at all is a separate question about the rule, not about a future count, and we test it on the audit set.

From the calibration set alone. With class probabilities and class-conditional thresholds, the calibration set can partly serve as its own audit. Whether an out-class compound receives the label *in-class* depends only on the threshold set from the in-class calibration compounds, so for this event the out-class calibration compounds are exchangeable with future out-class compounds and the audit construction applies to them. Whether an out-class compound *keeps* the label out-class is a coverage event, whose batch count follows the exact rank law used for the thresholds.[25] Each cell of a row is one of these two events or what remains after them, so computing both at level $1-\gamma/2$ bounds every cell of the row; a mixed batch adds an interval for the number of in-class compounds (ESI Section S4). Evaluating a class's own threshold on the compounds that set it, as is done when ordinary rates are computed on the calibration set, is not allowed; it makes the counts too narrow (ESI Section S5).

2.6 Experiments

Simulations. Eight designs are described in ESI Section S8:

1. audit-set intervals, with 10^4 replicates, each drawing its own data distribution;
2. finite libraries: 20 fixed libraries, each split at random 1000 times, which gives exchangeable data that are not independent;
3. threshold correction, with absolute Cauchy scores, $n =$ 50–500, $m = 100$ and 10^6 draws;
4. expected cost of the swept settings, computed exactly on the population, and selection: 4000 replicates of choosing among 100 settings under an assay capacity, with and without a Bonferroni correction;
5. scaffold dependence: compounds in clusters that share a shift of the predicted probability, with the future batch

drawn from the calibration clusters or from new clusters;

6. tied scores: a cluster of identical scores at a threshold, with ties kept or broken at random;

7. calibration-only intervals: 3000 replicates per calibration size, compared with spending the same labelling budget half on calibration and half on an audit set.

8. a two-level draw, in which a rate is drawn and then the calibration, audit and future counts are drawn given it, used to separate a statement about the long-run rate from a statement about the realized count of one batch.

Data. Molecules were standardized with RDKit[45], which keeps the largest fragment, neutralizes it and writes canonical SMILES; duplicates were then merged. One model was used throughout: Morgan fingerprints[46] (radius 2, 2048 bits) and CatBoost gradient-boosted trees[47] (500 trees, depth 6, learning rate 0.1) with fixed hyperparameters. The experiments concern the coverage settings and their audited consequences for a given model, not the choice of model. For the twelve Tox21 endpoints,[48,49] each endpoint was split 50/25/25 into training, calibration and audit sets: 50 times at random and 20 times by Bemis–Murcko scaffold.[50] Under the scaffold splits we also drew pilot samples of 200, 500 and 1000 compounds at random from the audit scaffolds, used each as the calibration set, and screened the remaining compounds. For AqSolDB,[51] compounds with $\log S \geq -4$ were labelled soluble. Ten stratified 70/15/15 splits were made, the model was trained on all training compounds, and calibration and audit were restricted to $\log P > 3.5$. Solubility most often limits lipophilic compounds, so a screen of this subpopulation is a realistic use; it is also a subpopulation on which the model's scores are distributed differently from the training data. For CALiSol-23,[52] a compilation of measured conductivities of lithium-salt electrolytes, each formulation (salt, concentration and solvent fractions) was labelled good if its conductivity at 20–30 °C was at least 8 mS cm^{-1}. The same model was trained on all temperatures of the training formulations, with temperature, solvent fractions and fraction-weighted solvent descriptors as inputs, and scored each formulation at 25 °C. Four designs, 50 repeats each, drew calibration and audit sets from random formulations, from held-out salt–solvent systems, or, for a new salt, from the other salts or from a pilot sample of the new salt. To check the audit intervals on real data, each audit set was split at random, 20 times, into an audit half and a stand-in future batch. A random halving makes the two halves exchangeable by construction, so these checks verify that the construction and its implementation are right and cannot fail for any model, representation or label noise; they are not evidence that an audit prices a batch assembled some other way. The designs that test that are the block holdouts and the declared round, where the batch is a distinct set of molecules. For the DNA-encoded library,[53] a sample of one million molecules was taken from the released table and a molecule was labelled in-class if it carried the library's own binding call for soluble epoxide hydrolase, a rate of 0.74%; the analysis runs on a further random subsample of 200 000 of those molecules, at a rate of 0.72%. Each molecule is built from one building block at each of three positions, so designs were defined over blocks rather than molecules: assignment at random, and assignment of whole blocks at one of the three positions, so that a batch molecule carries a block the training molecules do not. A declared round, 30% of it built on first-position blocks absent from training, was then prepared for in three ways that differ only in where the calibration and audit sets draw their novel chemistry. Counts for this library are for the realized batch of each repeat, tens of thousands of candidates, because at a sub-percent rate a batch of 1000 carries too few in-class molecules for a confirming setting to be bounded away from zero; the scripts additionally scale each rate to a nominal 10^5 so repeats can be compared, and those scaled figures are not the ones quoted here. The thresholds for this library were planned at $m = 1000$ rather than at the realized batch size, which makes the deployed rule more conservative than the request rather than less; ESI Section S14 gives the consequence, and it should be repaired before the library results are relied on. A further division of the electrolyte formulations by the paper each measurement came from, rather than at random, is described in ESI Section S10. Details, counts and seeds are in ESI Section S9, and for the library in ESI Section S14.

3 Results

3.1 Predicted counts hold at the stated certainty

The construction is classical — tolerance limits for the number of defectives in a future lot, applied to each cell of the outcome table — so the first question is only whether it is implemented correctly. It is. Under random splitting, which imposes exchangeability by construction and therefore tests the implementation rather than the assumption, realized counts in a batch of 1000 fell inside their intervals in 95.7–98.9% of replicates for every outcome and every audit size, and the same held for fixed libraries split at random, a design that is exchangeable but not independent (ESI Section S8). Two shortcuts a team might reach for instead each omit one of the two sources of variation and fail where the omitted one dominates: binomial quantiles at the audited rate, which treat that rate as known, covered as little as 44% of batches; m times a Clopper–Pearson interval for the rate, which ignores the draw of the batch, fell to 86.6–94.3% once the audit reached 1000 and is therefore unsuitable for a single plate. Bayesian posterior predictive intervals were 1.3–5.9% narrower and close to nominal here but carry no guarantee, falling to 94.8% in a two-level design. The intervals are also far from trivial: at the chosen setting of each scenario screen below they spanned 16–109 counts for a batch of 1000, 9 to 63 times narrower than $[0, m]$, and 36–204 after the joint correction over the sweep.

On real data the audit intervals behaved the same way (Table 3). Splitting each Tox21 audit set into an audit half and a stand-in future batch, the intervals covered 95.6–99.0% of the realized counts, for every outcome. Per endpoint, coverage was 95.8–97.1% under random splits. Those are pointwise figures, one outcome at a time, and the object a team is handed is the whole table: all eight lines of one setting — the six cells and the two aggregates — held together in 82.7% of random splits and 81.9% of scaffold splits. Eight pointwise 95% intervals do not make a 95% table, and a team that needs the table jointly must divide the level by the number of lines

Table 3: Real-data check of audit intervals (nominal 95%). Each Tox21 audit set (1405 to 1755 compounds) was split 20 times at random into an audit half and a stand-in future batch of the same size; entries are the percentage of realized future counts inside the interval. Random: 600 splits, 12,000 checks per outcome. Scaffold: 240 splits calibrated on other scaffolds, 4800 checks per outcome.

Outcome	random	scaffold
correct rejection	96.2	95.6
false lead	95.9	96.2
deferred out-class	95.7	95.6
confident hit	96.6	96.5
missed in-class	99.0	99.0
deferred in-class	96.9	96.4
single label	95.8	95.8
deferred	95.8	95.8

it reads. It was 95.0–97.5% under scaffold splits, in which the model had been calibrated on other scaffolds: the audit described the consequences correctly because it came from the compounds being screened. On AqSolDB, over all 132–144 settings of the sweep and all outcomes, the pointwise intervals covered 96.1% of 216,960 realized counts. With the Bonferroni correction over the 144 requests, every interval of every setting covered its count together in 98.5% of 200 audit–future splits.

Thresholds for small calibration sets. With the standard threshold, coverage over a batch of 100 compounds fell below the requested 90% in 35–42% of calibration draws for $n = 50$ to 500, against a requested tolerance of $\delta = 10\%$; the batch correction brought this to 4.7–9.9%, matching theory to within Monte Carlo error (ESI Table S1). The batch-free correction is identical for $n \leq 100$ but too permissive beyond it, falling short in 11.8% of draws at $n = 150$ and 21.2% at $n = 500$, so the batch form should be used whenever batch performance matters; the one-sided DKW correction was valid but far more conservative. The same holds on real data: per-class audit coverage fell below 90% in 55 of 600 random Tox21 splits for out-class compounds and 45 for in-class compounds, against 56.2 and 42.3 expected from the exact law for the thresholds used (ESI Table S13), where the batch-free correction fell short in 17% and 13% of splits and the DKW correction in 0.2% and 1.0%.

3.2 What accepting the default setting costs

A sweep is only worth tracing if the setting a team would otherwise accept is a poor one. On the population it is: the default request of 90% for both classes, chosen without the costs in view, cost up to 3.1 times the cost-optimal two-threshold rule when a miss was ten times a false lead, and up to 32 times when an assay was nearly free (ESI Table S4).[36] The swept settings contained that optimum in 46 of 60 cost regimes. None of this claims to beat a rule built for a known cost ratio — where the ratio is known and the probabilities can be recalibrated, such a rule is the thing to use, and on the solubility screen the swept setting is indistinguishable from one built on isotonically recalibrated probabilities (ESI Section S7). The case for the sweep is that it needs no ratio and returns audited counts for whatever is chosen.

Choosing the setting from audited outcomes needs the correction for selection. Among 100 swept settings we took the one with the fewest audited errors whose upper limit on deferrals fitted a capacity of 15% of the batch. With the Bonferroni correction, realized deferrals exceeded capacity in 0.03% of replicates; without it, in 0.47%. Both are below the nominal 5%, because neighbouring settings have strongly correlated counts, but only the corrected choice carries a guarantee, and it cost 52.5 rather than 43.5 errors per 1000 compounds.

3.3 Sizing a round

The same intervals run backwards answer the question a campaign asks first: given a target number of hits, how many candidates must the batch hold? The smallest batch whose lower bound on confident hits reaches the target answers it. The interval is valid for each batch size fixed in advance and the search makes the size a function of the audited count, so what follows should be read as the size that count implies rather than as a size carrying its own guarantee; ESI Section S3 gives the convention, which is one-sided, and the qualifications.

The answer differs from the one an expectation gives, and most where screens operate (Fig. **??**). Dividing the target by the observed rate, which is what a sum of calibrated probabilities amounts to,[20] asks for 1000 candidates to obtain 20 hits at a rate of 2%. The audited count implies 2056 on an audit of 384 and 1687 on an audit of 1000; at a rate of 50% the expectation asks for 40 and both audits give 52. Three quantities are stacked in that gap and separating them is what makes it planable. Were the rate known exactly, 1390 candidates would still be needed, because which candidates fall in this particular batch is not something measurement can settle. Knowing the rate from an audit of 384 raises it to 2056, and from an audit of 49 to 19 091. The first is the risk the campaign carries; only the second and third are for sale. A team planning on the expectation omits all three.

At these rates the requirement is better read as a distribution than as a number: an audit of 49 at a rate of 2% sees no in-class compound with probability 0.37, in which case no batch size reaches the target at all. Enlarging the audit from 49 to 384 removes that possibility and roughly halves the typical requirement. On this accounting the audit is not an overhead on the campaign but the measurement with the largest effect on what the campaign costs, and its effect is largest exactly where a screen's hit rate is low.

How much is worth buying can be settled before any compound is measured. Accepting a width r times the narrowest any audit could achieve requires $n = m/(r^2 - 1)$ audit candidates, a ratio independent of the rate and of the batch size: an audit the size of the batch lands a factor $\sqrt{2}$ above that floor, one of $2.3m$ within 20% of it. The floor itself depends on the rate and at primary-screen rates it is wide — for a batch of 1000 at 2% it is ± 9 on an expected 20 — so a team should read the floor first and enlarge the batch, not the audit, when the floor is what binds. Correlated series change the two statements differently: the floor is a width and inflates by the square root of the design effect, while the ratio rule needs no inflation, since the design effect multiplies both variance terms and cancels from r.[54,55] Re-estimating the rate between rounds costs a seventh to a quarter of a batch if the number of looks is declared in advance, and is

not covered at all if a campaign stops when the hits so far look sufficient; ESI Section S3 states both conditions and the anytime-valid alternative.

3.4 What a measurement budget costs: a solubility screen

After standardization and de-duplication, 9254 AqSolDB compounds remained, 2467 of them with $\log P > 3.5$. In each of ten splits, the lipophilic calibration and audit sets held about 370 compounds each, 22% of them soluble ($\log S \geq -4$). Audit AUROC was 0.91–0.93 overall and 0.82–0.88 within the lipophilic subset.

What the labels can carry. The solubility label cuts a continuous property measured across sources, so it is noisiest where it matters, at the cut: about 8% of the screened compounds near it would be labelled the other way by another laboratory, which is the size of the requests being swept. Rerunning the whole analysis on the same scores with only the labels changed shows what that supports and what it does not. The resolution of the sweep survives, and so does the enrichment, at a fold between 1.51 and 3.61 and significant after correction in every one of a hundred perturbed evaluations — so an enrichment of about twice the base rate is safe and the two digits of 2.4 are not. The scenario does not survive: the number of certifiable settings falls from a median 26 to 18 under random flips and to 2 and 0 under the two adversarial variants, and the cheapest certifiable budget rises by a factor of 2.2 to 2.8. About a quarter of the Pareto front is an artefact of the labels. One detail runs against the mechanism of Section 3.11 and is reported there: flipping the compounds the score ranks most extremely *flatters* the screen, because noise correlated with the score makes the labels agree with the model, so a proxy that agrees with the model conceals its own failure rather than revealing it (ESI Section S13).

The price of a budget. Cutting the measurement budget from 450 to 250 per 1000 lipophilic compounds cost 73 additional false leads, with all requirements guaranteed jointly at 95%, and the cheapest budget any setting can be guaranteed to meet is 173 measurements per 1000, at which exactly one setting remains certifiable and its guaranteed false leads rise to 335. At this audit size the guarantee itself cost about 110 measurements per 1000: the chosen setting expects 286 and is bounded at 395. Those figures are one split, and the spread across the ten is wide enough to matter: the number of certifiable settings runs from 11 to 65 and the cheapest certifiable budget from 29 to 309 per 1000, with a median of 134. The setting chosen here is itself certifiable in only four of the ten splits, so it is this split's choice and not a recommendation. The rest of this section shows how these numbers follow from the map.

The decision space. Sweeping (α_0, α_1) over 12 values each, with $\delta = 0.10$, gave 132–144 distinct settings per split. Fig. 2 (left) shows their audited consequences for one split: one model, one endpoint, and outcomes ranging from 10 missed soluble compounds per 1000 with a third of the batch sent to measurement, to almost no measurements at the price of about 50 missed soluble compounds. Of the settings, 74–120 per split were Pareto-optimal for missed soluble compounds, false leads and measurements (ESI Table S11). A front that large invites the question whether it is a decision surface or a restatement of the grid, and the answer depends on the objective: minimising missed compounds selected the same request in eight of the ten splits and one two grid steps away in the other two, while minimising deferrals wandered over eight distinct requests, all within one corner of the grid. The front is therefore stable where the objective is steep and loose where it is flat, and a team should read the chosen setting's audited counts rather than its coordinates. So the front is the menu, and the team's constraints make the choice. Table 4 gives the audited consequences of the two extremes, with intervals for a batch of 1000. Setting A misses 10 soluble compounds per 1000 and sends 335 to measurement; setting B sends none to measurement and misses 52. Because A and B were picked by looking at the audit, the intervals that hold for them are the Bonferroni ones of Table 4, $[0, 49]$ and $[15, 114]$ missed, rather than the pointwise $[2, 29]$ and $[29, 84]$ that the same table reports for a setting fixed in advance. Over the ten splits the corresponding audited values were 0–20 missed and 203–522 deferred for setting A, and 26–61 missed and 0–42 deferred for setting B.

Scenarios. Two scenarios illustrate the choice (Fig. 2, centre and right). In both, a batch of 1000 lipophilic compounds must yield at least 60 confirmed soluble compounds and miss at most 90, and the objective is the fewest expected false leads; the first allows at most 450 compounds to be sent to measurement, the second at most 250. The limits were set after inspecting the map, which the joint bounds allow: every constraint is checked against a one-sided audit bound with the Bonferroni correction over the 144 requests and four bounded outcomes, so all of them hold together with probability at least 95%. Under the first budget 26 of the 144 requests were feasible, and the chosen setting expects 115 false leads, 286 measurements and 23 missed soluble compounds per 1000, bounded jointly at 395 measurements, 75 misses and at least 66 confirmed. Tightening the budget to 250 left 9 feasible requests and a setting expecting 188 false leads, so the tighter budget costs 73 additional false leads per 1000. The price of a budget can be read off directly over a range: 221 false leads at 200 measurements per 1000, 188 at 250–300, 133 at 350, and beyond 375 the budget no longer binds. Four settings remain certifiable at 200 and one at 173, below which nothing meets all requirements with the labels in hand.

Those figures are one split, and the spread across the ten is wide: the number of certifiable settings runs from 11 to 65 and the cheapest certifiable budget from 29 to 309 per 1000, with a median of 134. The setting chosen here is itself certifiable in only four of the ten splits, so it is this split's choice and not a recommendation. Two further prices are worth naming. The joint correction costs reach as well as width — the share of the grid satisfying the scenario's limits falls from 0.2–0.6 read pointwise to 0.0–0.1 read together — and the gap between expected and bounded counts, 286 against 395 measurements, is the price of the guarantee at this audit size, which labelling buys down: half the audit would give a bound of 432, twice the audit 372 and four times 358. The compounds advanced without a measurement were 54% soluble against a prevalence of 22%, an enrichment of 2.4

times on 51 of 95 compounds and significant after correction over the sweep, so the model earned the measurements it saved. No cost ratio is needed for any of this; the ratios for which each setting is expected-cost optimal are mapped in ESI Fig. S2.

3.5 One request, twelve sets of consequences: Tox21

After standardization and de-duplication, 7592 Tox21 compounds remained, with 5619–7017 labelled per endpoint and an in-class prevalence of 2.9–15.9% among labelled compounds (the labels are missing non-randomly, which the Limitations quantify). Here a confident hit is a compound confidently flagged as toxic, which a screen would set aside. Calibration sets contained a median of 44–228 in-class compounds per endpoint (ESI Table S13). Audit AUROC was 0.70–0.87 for random splits and 0.67–0.85 for scaffold splits. These endpoints also bound the decision-geometry reading of Section 2.4, which needs the probability of being in-class to increase with the score. It does not here: the in-class fraction decreases over at least one adjacent pair of score bins in nine of the twelve endpoints, nominally so for two at $z = 2.4$ and 2.6 — an uncorrected comparison that does not survive correction across the 108 bin pairs examined — and for three endpoints the fraction is nearly flat over the lower nine bins with the signal confined to the highest (ESI Table S10). The containment argument and the cost-optimal reading therefore do not apply to those endpoints, though the intervals themselves need no such condition. Before tie-breaking, a median of 3.3–4.3% of scores were tied, mostly from identical fingerprints.

The same request, very different consequences. All endpoints were calibrated with the same request, $\alpha = \delta = 0.10$ per class, using the batch correction. Yet the median fraction of compounds receiving a single label ranged from 13.6% (NR-PPAR-γ) to 62% (NR-AhR) (Fig. 3). Table 5 gives the audited consequences of the two extremes, for the split with the median single-label fraction. For NR-AhR, a batch of 1000 compounds yields 78 confident hits (audit interval 58–101) against 78 false leads (57–100). For NR-PPAR-γ, it yields 15 (6–26) against 89 (67–113), with 865 compounds deferred. Half of the toxic calls are correct for NR-AhR, one in seven for NR-PPAR-γ. The audit reveals these consequences before the screen is run. The fraction of in-class compounds among single-label "in-class" calls ranged from 13% to 58%, an enrichment of 3.0 to 6.5 times over prevalence. With the standard threshold the single-label fractions were larger, 28% and 72%, but that threshold's coverage of a class fell below the 90% target in 41.5% of the 600 endpoint–split pairs for in-class compounds and 47% for out-class compounds, and in 54–80% of splits per endpoint for at least one of the two.

3.6 Auditing a new salt: an electrolyte screen

CALiSol-23 holds 13,825 conductivity measurements of lithium-salt electrolytes from 27 publications.[52] Three source counts appear below and each belongs to a different pool: 27 in the released file, 26 after the exclusions described next, and 25 contributing formulations measured at room temperature, which is the pool every experiment here uses. After removing 523 rows without a conductivity and 23 rows from one source whose values lie about 100 times below all others for the same formulations (a probable unit error), 1495 formulations had been measured at room temperature, 275 of them (18%) at or above 8 mS cm^{-1}. The threshold is close to the conductivity of standard carbonate electrolytes; the ten room-temperature measurements of 1 M $LiPF_6$ in EC/EMC (3:7) in this dataset run from 7.85 to 10.21 mS cm^{-1}, so the reference electrolyte itself straddles the cut and a good formulation is one within reach of that standard rather than reliably above it. That spread is the price of a binary label on a continuous property measured across sources, and a quarter of the formulations (374 of 1495) sit within 20% of the cut. A threshold of 10 mS cm^{-1} would leave 102 good formulations in the pool, about 24 in a calibration set of 350, which is just above the attainability floor (Section 3.12) and would leave the request barely reachable. Measuring conductivity is cheap compared with building and cycling cells, so here a deferral means a measurement.

A new salt. The sharpest test holds out a whole salt. It is a weaker test than the name suggests: the salt held out is $LiPF_6$, which is the incumbent of commercial practice, supplies 606 of the 1495 room-temperature formulations, and anchors the class boundary, so what is tested is transfer to a well-represented salt rather than to novel chemistry; its share of good formulations, 25%, also differs from the salts calibrated on, so prevalence moves with the chemistry here as it does in the library screen. A genuinely novel salt is not available in this compilation, where several candidates contribute a single room-temperature formulation each. With a model trained and calibrated on the other salts and the 606 $LiPF_6$ formulations screened, coverage of the good class averaged 29% and fell below 90% in all 50 splits (Table 6; Fig. 4, right). The audit of the screened formulations showed what this meant for the screen: at the default setting it reported a median of 174 missed good electrolytes per 1000 formulations, about two-thirds of the good ones, and 59 confident hits. Calibrating instead on a random pilot of 40% of the $LiPF_6$ formulations, about 240 measurements, restored it for the good class (5 shortfalls against 3.3 expected) and left a residual excess for the poor class (8 against 4.0, one-sided Poisson-binomial $p = 0.04$), and the audit showed the gain: 10 missed good electrolytes and 159 confident hits per 1000, for 308 measurements. Formulations split at random over all salts met the request (2 and 4 shortfalls against 4.7 and 3.1 expected), and the sweep of coverage requests gave 108–132 distinct settings per split, trading measurements (0–230 per 1000) against false leads (0–150) and missed good electrolytes (0–30) (Fig. 4, left). Holding out whole salt–solvent systems also broke the coverage guarantee (20 and 22 shortfalls against 4.5 and 5.0).

The audit carries the guarantee. In all four designs, including the two in which the coverage guarantee failed, an audit set from the screened formulations reported the consequences correctly. Audit intervals for the default setting covered 95.5–100% of realized counts, pointwise intervals over the whole sweep 97.2–97.6%, and with the Bonferroni correction every setting and outcome was covered together in 98.4–99.2% of audit–future splits. A capacity rule chosen on

Table 4: Audited consequences of two settings of the lipophilic solubility screen (split 0), per 1000 compounds: realized rate per 1000 in the audit set (384 compounds; 10 per 1000 is 4 observed compounds), with the 95% prediction interval for a batch of 1000 and the interval with Bonferroni correction over the 144 requests of the sweep, which remains valid for one nominated outcome after choosing a setting by its audited outcomes. Holding all six lines of a chosen setting jointly would require the level divided by 144×6 instead; the scenario bounds quoted in the text are computed that way, over the four outcomes they constrain. A: the Pareto setting with fewest audited missed soluble compounds, ties broken by fewer deferrals, $(\alpha_0, \alpha_1) = (0.30, 0.075)$; B: the Pareto setting with fewest audited deferrals, ties broken by fewer misses, $(0.275, 0.25)$; $\delta = 0.10$.

	Setting A			Setting B		
	audit	95%	Bonferroni	audit	95%	Bonferroni
missed soluble	10	[2, 29]	[0, 49]	52	[29, 84]	[15, 114]
false lead	232	[183, 285]	[148, 330]	221	[174, 274]	[139, 319]
soluble confirmed	177	[134, 226]	[103, 269]	172	[129, 220]	[99, 263]
insoluble rejected	245	[195, 299]	[159, 345]	555	[495, 614]	[446, 660]
deferred, insoluble	299	[246, 356]	[206, 404]	0	[0, 11]	[0, 26]
deferred, soluble	36	[17, 64]	[7, 92]	0	[0, 11]	[0, 26]

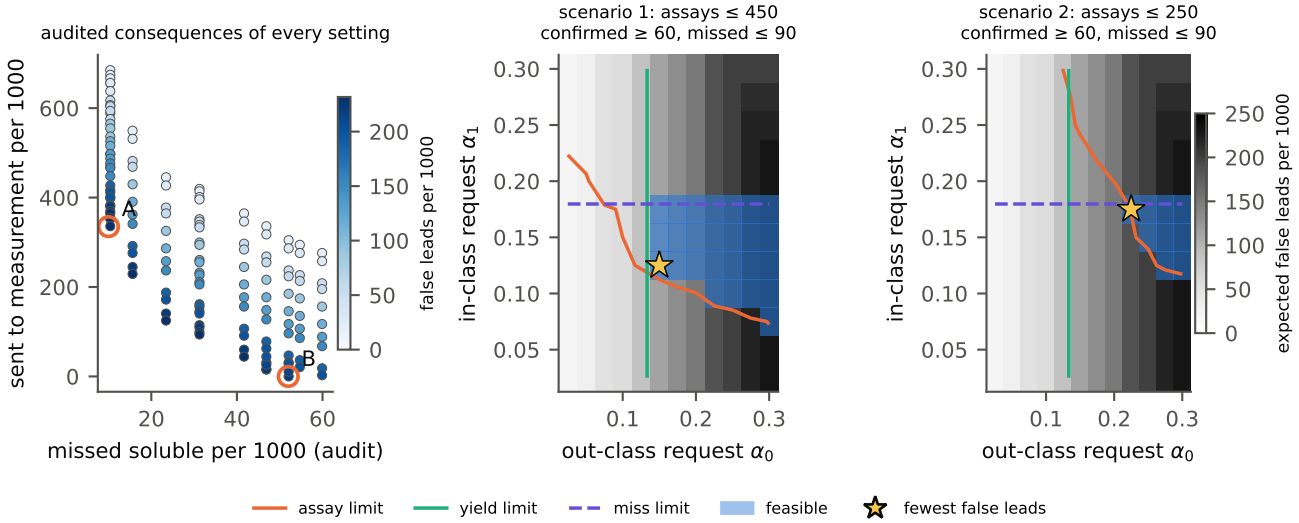


Figure 2: Navigating a solubility screen for lipophilic compounds (AqSolDB, $\log P > 3.5$; split 0) with one model. Left: each point is one setting of the coverage requests, placed by its audited missed soluble compounds and measurements, with false leads in colour; A and B are the settings of Table 4. Centre and right: two scenarios in the plane of the coverage requests, with the expected number of false leads (grey, the objective) for each request. Lines mark where the 95% joint bound of each constrained outcome equals its limit; blue cells satisfy all constraints with joint probability at least 95%, and the star is the feasible setting with the fewest expected false leads. The right panel tightens the assay budget from 450 to 250 per 1000.

one audit half, as a planner would choose it, never exceeded capacity on the other half, in 980–1000 checks per design. The intervals were also informative: for the scenario setting the 95% intervals were 12 to 63 times narrower than the trivial $[0, 1000]$, and 7 to 28 times narrower jointly over the sweep. The advanced formulations were 71% good against the 15% prevalence of the audit set of the split these maps use, which is lower than the 18% of the room-temperature compilation as a whole, an enrichment of 4.7 times (one-sided Fisher exact test on the audit set, $p = 2 \times 10^{-37}$, Bonferroni over 144 settings). Scenario maps for the electrolyte screen behave as for solubility (ESI Fig. S3). Intervals from the calibration set alone covered 99.7% of counts for random formulations; for a new salt they rely on a calibration that does not match, and a separate audit is the tool to use.

Per plate. A self-driving platform measures plates rather than batches of 1000. For the pilot-calibrated rule, an audit of 364 formulations gives, for one 96-well plate, 30 expected measurements with a 95% interval of 20–40; scaled Clopper–Pearson would give 25–35, too narrow at this ratio of audit to batch size (Section 3.1). The one-sided bound of 38 measurements per plate holds with probability 95% for a single plate, and a bound of 43 holds for all ten plates of a campaign together (level 0.5% per plate).

3.7 An audit and a batch that are not two halves of one set

Every check of the audit intervals reported so far divides one labelled pool at random, which makes the two halves exchangeable by construction and establishes that the construction is right rather than that it survives a real separation. The electrolyte compilation lets the harder question be asked, because it records which paper each measurement came from. We therefore divided the formulations by source instead of at random, so that no paper contributes to both the audit and the batch, holding the model and the calibration set fixed across the divisions and halving nothing. The protocol is the one used elsewhere in this paper: formulations are assigned to

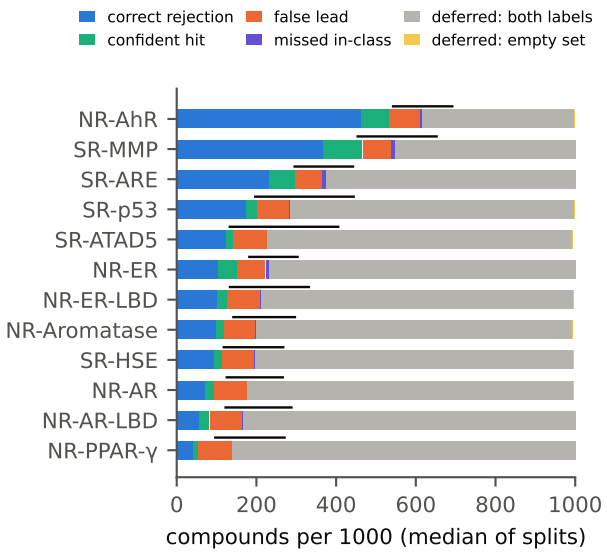


Figure 3: Twelve Tox21 endpoints calibrated to the same request ($\alpha = \delta = 0.10$ per class, batch correction) sort a screening batch very differently. Bars: median audited counts per 1000 compounds over 50 random splits, ordered by single-label fraction (medians need not sum to 1000); no empty sets occurred. Black lines: 5–95% range of the single-label count over splits.

Table 5: Audited consequences per 1000 compounds for the endpoints with the highest and lowest single-label fraction (batch correction), for the random split with the median single-label fraction (25th of 50). Entries: realized rate per 1000 in the audit set, with the 95% audit interval for a batch of 1000 in brackets. Calibration sets: 1385 out-class compounds and 197 in-class compounds (NR-AhR), 1526 and 37 (NR-PPAR-γ).

True label	Outcome	NR-AhR	NR-PPAR-γ
out-class	correct rejection	456 [417, 496]	32 [19, 48]
out-class	false lead	78 [57, 100]	89 [67, 113]
out-class	deferred	342 [305, 380]	850 [820, 877]
in-class	confident hit	78 [58, 101]	15 [6, 26]
in-class	missed in-class	7 [2, 15]	0 [0, 3]
in-class	deferred	39 [24, 55]	15 [6, 26]

training, calibration and the remainder, the model is fitted on every temperature row of the training formulations and is therefore disjoint from the set that sets its thresholds, and the remainder is divided three ways.

The intervals do not survive it. Over ten repeats, random halving covered 0.97 ± 0.04 of the consequence lines, as it must; partitioning the papers at random and letting each formulation follow its own paper covered 0.70 ± 0.12; and taking one paper's formulations as the batch and the rest as the audit, over the five papers large enough to serve within the remainder, covered 0.77 ± 0.11. The inference here is about papers, so the paper is the unit and not the repeat: the five per-paper means run from 0.58 to 0.90, and pooling the 41 runs instead would quote ± 0.05 and hide that spread. A division by provenance therefore costs roughly a quarter of the coverage a random halving delivers, and nothing a campaign can compute beforehand announces it. Area under the curve on the batch ran from 0.81 to 1.00 across these divisions. The enrichment of the confident calls, taken over the pool's base rate, which is the rate a campaign knows, sits between 3.3 and 4.9 in all three divisions and does not track coverage within any of them (Spearman $+0.01$ by source, -0.19 one source out). The one quantity that does track coverage is the ratio of the audit's in-class rate to the batch's, and that ratio needs the batch's labels, which is what the audit was to replace.

What the designs do not do is separate the causes. The median ratio of audit to batch in-class rate is 1.02 halved, 0.67 by source and 1.88 one source out, and the per-repeat ratios span 0.3 to 2.2 and 0.4 to 8.1 in the two source divisions. So dividing by provenance moves the in-class rate as well as the mapping from score to outcome, and the coverage loss is attributable to the two together. That is itself the practical point — a paper is a population in both senses — but it means this experiment shows that provenance is enough to void the intervals, not that the mapping alone is.

A pilot of the right population, divided for real. The remedy this suggests is the one the library screen also points to: when a round reaches into a population the audit does not hold, buy a pilot of that population. Every arrangement reported so far that recovers the nominal level divides one pool at random, which cannot fail, so none of them tests that advice. The compilation lets it be tested, because a single paper's formulations can be divided in ways that are not random. Holding the thresholds on the other sources throughout, so that the rule is calibrated off-population exactly as a campaign would have it, we held out one source as the screened population and varied only where the audit came from: from the other sources, as a team would do by habit; or from the held-out source itself, divided by concentration band, by solvent system so that no system is shared, or simply as the paper reported its formulations, first half against second.

The advice holds. Taking the source as the unit, an audit drawn from the other sources covered 0.55 ± 0.07 of the consequence lines over eight sources. A pilot of the screened source covered 0.92 ± 0.17 when the division was by concentration, 0.87 ± 0.21 by solvent system and 0.86 ± 0.12 as reported; on the sources that support both arrangements the paired gain over auditing off-source is $+0.37 \pm 0.21$, $+0.37 \pm 0.26$ and $+0.34 \pm 0.17$. A random division of the same source covered 0.99 ± 0.02, which is the ceiling the construction guarantees and which these divisions do not quite reach: sampling the right population recovers most of what provenance costs, and not all of it.

Two details are worth keeping. The divisions by concentration and by system leave the audit and batch in-class rates close, at median ratios of 1.04 and 1.09, so this is not a prevalence effect; the off-source arm sits at 0.56. And the division as reported, which is not chemically meaningful but merely not random, leaves the rates badly mismatched at a median ratio of 1.81 and still recovers to 0.86. Sampling the population the batch comes from therefore helps even where the prevalence does not match, which is the asymmetry the construction predicts: what the audit must share with the batch is the mapping from score to outcome, and a prevalence gap it can report.

Two things follow. The first is a limit on this work's own evidence: the coverages quoted for solubility, toxicity and electrolytes under random splitting are verification, and the only divisions in this paper that test the intervals against a genuinely distinct batch — these, and the library's block

Table 6: Electrolyte screen (CALiSol-23, conductivity $\geq 8\,\mathrm{mS\,cm^{-1}}$ at room temperature), 50 splits per design. Shortfalls: splits in which audit-set coverage of the class fell below 90%, observed/expected under the exact law. Audit coverage: percentage of realized counts in the stand-in future batch covered by 95% audit intervals, for the default setting (range over outcomes), for every setting of the sweep (pointwise), and for all settings and outcomes together with the Bonferroni correction over 144 requests.

		Shortfalls		Audit coverage (%)		
Calibration and audit drawn from	AUROC	poor	good	default	sweep	Bonferroni
random formulations	0.98	2/4.7	4/3.1	95.5–99.4	97.2	98.4
new salt–solvent systems	0.91	20/4.5	22/5.0	96.6–98.9	97.6	99.2
new salt, calibrated on others	0.92	0/4.5	50/1.9	95.9–100.0	97.3	98.9
new salt, pilot calibration	0.92	8/4.0	5/3.3	96.0–99.2	97.2	98.5

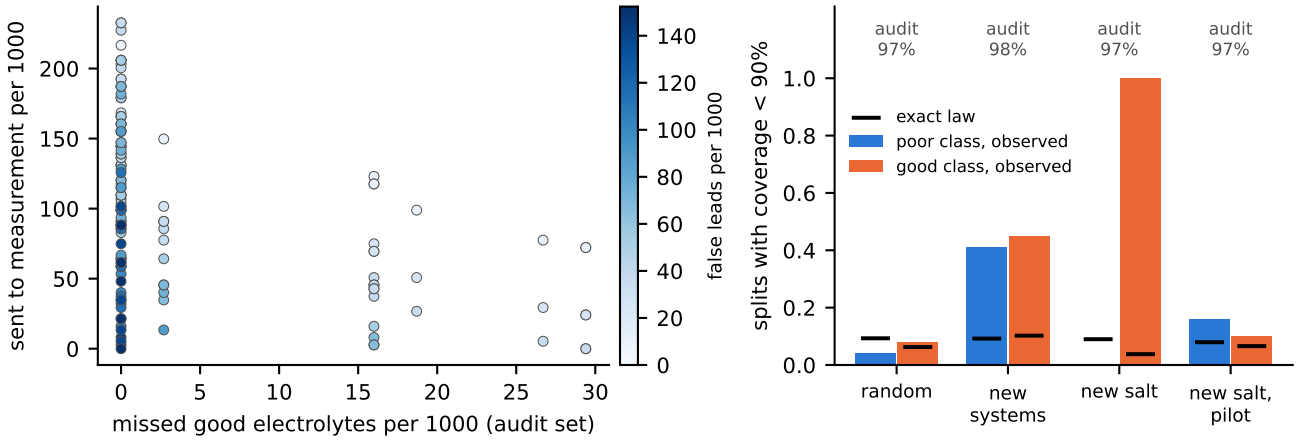


Figure 4: Electrolyte screen (CALiSol-23). Left: each point is one setting of the coverage requests for one random split, placed by its audited missed good electrolytes and measurements per 1000 formulations, with false leads in colour. Right: fraction of 50 splits in which coverage of each class fell below 90%, observed (bars) and expected under the exact law (black marks), for calibration on random formulations, on other salt–solvent systems, on other salts, and on a pilot sample of the new salt; above each design, the pointwise coverage of the audit intervals over the sweep.

holdouts — break them. The second is a limit on pooled compilations, which is where much of the data for materials screening comes from. An audit drawn from the library to be screened is not sufficient if the batch will be measured by someone else, on another instrument, with another calibration: provenance is part of what must be exchangeable, and a compilation of 25 papers is 25 populations. The instruction to audit the population you will screen therefore extends to the laboratory that will screen it.

3.8 Building the audit for the round it prices: a DNA-encoded library

The electrolyte screen showed that an audit drawn from the screened population bounds the consequences where a calibration set from elsewhere does not. A campaign poses a harder version, because the next round often introduces chemistry that has not been made. What the audit must then contain is a question about its design, and a DNA-encoded library answers it cleanly: the chemistry is built from a catalogue of building blocks, so a round can be declared in those terms.

What a structural split costs. Assigning molecules to training, calibration, audit and batch at random keeps all four exchangeable, and the intervals covered 53 of 54 consequence lines over nine repeats. Withholding instead every molecule carrying a given block at one position, so that the batch is built on chemistry absent from the calibration and audit sets, breaks them. Taking the repeat as the unit — the six lines of one repeat are cells of a single table on a single batch, not independent checks — coverage was 0.48 ± 0.19, 0.37 ± 0.14 and 0.35 ± 0.12 for the third, second and first positions, against 0.98 ± 0.04 for random assignment. Over the same runs the model's ranking fell only from 0.960 to 0.892 by area under the curve. This is the split a project adopts to guard against memorisation, and passing it licenses nothing quantitative about the next batch.

Nothing a campaign can compute beforehand orders these failures. Area under the curve is governed by the ordering of the negatives at a sub-percent rate and is invariant to prevalence by construction, so it cannot announce a prevalence change. The enrichment of the confident calls, which a screening group reports as a matter of course, is the audit's own in-class rate re-expressed — it divides the precision of those calls by exactly that rate — and within the first-position repeats it does not order them by coverage: its largest value belongs to a repeat covering three of six lines, while the repeat covering none sits mid-range. The audit's rate against the library's is a cheap partial warning rather than a test: pooled over the thirty block-holdout repeats it predicts coverage only moderately, and within the first position not at all (ESI Section S14). What does predict coverage is the audit's rate against the *batch's*, which needs the batch's labels.

These designs move the in-class rate of the audit and of the batch independently, between 0.2% and 2.5%, because

in-class molecules are concentrated in a few blocks. So the losses are a change in prevalence as well as a change in the mapping from score to rate, and the designs do not separate the two; the arm that holds matches the batch's prevalence by construction. Dealing blocks to the four roles in order of their in-class count would separate them while keeping the design block-disjoint, and we have not run it.

A declared round. A round is declared in which 30% of the batch is built on first-position blocks the model has never seen, and three preparations are compared, each supplying its own calibration and audit sets and all three sharing one model per repeat: calibration and audit from familiar blocks only, as a team would do by habit; from different new blocks but at the batch's own 30% share, which is where a team stands that has measured some new chemistry and is about to make different new chemistry; and from the batch's own new blocks, which is a pilot of the chemistry the round will contain.

Reproducing the batch's composition is worth nothing. Familiar blocks covered 0.63 ± 0.13 of the consequence lines and the batch's proportion of new chemistry from other blocks covered 0.63 ± 0.10; because those two arms share a batch and nearly nested audits their difference can be taken per repeat, and it is 0.000 ± 0.070, which is the strongest form of the claim. Neither comes near the request. Sampling the batch's own blocks covered 0.97 ± 0.04 and was the only arrangement to cover all six lines in a majority of repeats, eight of ten against one and none — but that arm splits a fixed pool at the molecule level, which makes its calibration, audit and batch exchangeable by construction, so its 0.97 is a positive control rather than evidence that a pilot repairs a campaign. Its narrower interval is likewise a property of its audit being 43% larger and not of the arrangement (ESI Section S14).

What distinguishes the three is therefore not how much new chemistry the audit holds but whether it holds *that* chemistry. The first is cheap and can be arranged before synthesis; the second requires measuring a sample of the new blocks, which is an experiment. Both arrangements that do not sample the batch's own chemistry fail, and none short of sampling it was found to work.

3.9 The batch the model chose

Every division so far has been random or structural. A campaign is neither: a planner reads the scores and commits the top of the ranking, which is the one arrangement this construction is not expected to price. We tested it on the library, holding the model fixed so that only the selection moves. A pool of 70 000 candidates was scored, the planner committed the highest-ranked m of those remaining in each round, and the same committed batch was priced three ways: with a random sample of the candidates still in the pool, which is the audit a team would reserve; with a random slice of the committed batch itself, measured and used to price the remainder; and with the same selection rule applied to a pool of 20 000 reserved from the planner. Two round sizes were run, $m = 8000$ and $m = 1500$.

The audit reserved from the library does not price a plate the model chose, and the reason is visible before any interval is computed. In the first round at $m = 1500$ the committed batch had an in-class rate of 0.269 against 0.0071 for the pool it was drawn from, a factor of 38; what remained of the pool fell to 0.0013. Across the twelve rounds of the two campaigns the library audit could not even be formed in nine, because a sample of the leftover pool contained no in-class candidate at all, and in the three rounds where it could be formed it covered 2, 2 and 4 of the six consequence lines. A quantity that cannot be estimated is a safer failure than one estimated wrongly, but neither is a guarantee.

Both repairs work, and the first is exact. A random slice of the committed batch covered all six lines in every one of the nine rounds in which it could be formed, which is what the construction requires: a random division of a committed set is exchangeable with the rest of that set however the set was chosen, so the selection drops out. Applying the selection rule to a reserved pool also worked, covering all six lines in eight of ten rounds and four of six in the other two; it is an approximation rather than an identity, since the reserved pool and the screening pool are two draws rather than one, and it costs a pool of high-scoring candidates the planner may not commit. The three rounds in which the batch was rich enough for the comparison to carry weight — batch rates of 0.269, 0.055 and 0.020 — show the pattern most clearly: library audit unavailable, 2 of 6 and 2 of 6; batch slice 6, 6 and 6; reserved selection 6, 4 and 6.

The second finding is about the sequence rather than the round, and it is the one that bounds what any audit can do for a campaign. One round exhausts the library. Committing 1500 of 70 000 candidates, 2% of the pool, took the in-class rate of the remainder from 0.0071 to 0.0013, and committing 8000 took it to 0.0005; by the fifth round of the smaller campaign no arrangement could be formed at all, because nothing in-class was left to sample. So a round's audit does not price the round after it even with the model frozen, the rule fixed and the selection unchanged, and not because the mapping from score to rate has moved: the population underneath has been removed. An audit is a measurement of the library as it stands, and at a sub-percent in-class rate a single round changes the library it measured. The instruction that follows is to reserve a slice of every committed round rather than to amortize one audit over a campaign, which is cheap, and to read the falling rate of the remaining pool as the campaign's own statement about how much is left to find.

3.10 Checking exchangeability before the batch is made

The preceding result is a design instruction, and a team following it would like to know whether it has succeeded. Exchangeability binds the calibration set, the audit set and the batch, and all three are sets of structures that exist before any of them is assayed, so the question can be put to the structures alone. We ask it as a two-sample problem: a classifier is trained to tell audit compounds from batch compounds using fingerprints and no labels. That is the pair the predicted counts depend on, since the audit prices whatever rule it is handed and asks only to be exchangeable with the batch; the calibration set may be compared with the batch in the same way when the coverage guarantee is also wanted. Under exchangeability the two are draws from one distribution, no classifier can separate them, and its area under the curve sits at one half.

Run across the four split designs, the test separated every arrangement that later failed and did not separate the one that held (Table 7); what it can refute is bounded by the representation, as the end of this section and the Limitations set out. It sits at 0.498 for random assignment, which covered 98% of consequence lines, and at 0.985, 0.989 and 0.997 for the three block holdouts, which covered 48, 37 and 35%; comparing the calibration set with the batch instead gives 0.501, 0.987, 0.990 and 0.999, so on these designs either pair would have raised the alarm. Beside it we report the share of batch molecules whose block at each position also occurs in the comparison set, which is the quantity the design controls directly. That share is one everywhere under random assignment and falls to zero at exactly the withheld position in each holdout, so the two statistics together give a verdict and name the axis responsible.

Those four designs are the easy case, and the table should be read as a check that the test fires rather than as a measure of its sensitivity: a withheld block catalogue is a substructure a fingerprint represents directly, so a classifier must find it, and the block-share column gives the same verdict by inspection. The question that matters is whether the test sees the three preparations of the declared round, which differ not in how much novel chemistry the batch carries — all three hold 70% of their batch's first-position blocks, or all of them for the pilot — but in where the audit draws its own. Over two repeats the audit-against-batch area under the curve was 0.658 for the naive preparation and 0.663 for the mirror, against 0.502 for the pilot; over the ten repeats of the coverage experiment the first two cover 38 of 60 consequence lines each and the pilot 58 of 60. The permutation null, obtained by shuffling the audit and batch labels and refitting at the same cap as the statistic itself, averaged 0.500 over 120 replicates and spanned 0.491 to 0.511, so the two preparations that fail are separated from the null by a wide margin and the one that holds lies inside it. The test therefore distinguishes arrangements that a fingerprint comparison of composition cannot, which is what the workflow step requires of it.

The test is one-sided and should be read that way. Separation is strong evidence against exchangeability, because two samples from one distribution cannot be told apart. Failure to separate is weak evidence for it, since a classifier may simply lack the features to see a difference that matters; the fingerprints used here cannot distinguish stereoisomers, for instance, whose in-class rates in this library differ severalfold. A screen that passes is therefore not certified, but a screen that fails is refuted, which is enough to be worth running: it costs no measurements, uses structures already known, and can be run before the plate is committed. The clearest case it cannot see is the provenance division of Section 3.7: which paper a measurement came from is not a property of the formulation, so no test on structures can detect a division that costs roughly a quarter of the coverage. At the first position of this library, 60 of the 294 blocks form 30 pairs that are stereoisomers, indistinguishable under the achiral fingerprint used here and differing severalfold in in-class rate; that is a fifth of the catalogue carrying most of the binders, invisible to the test by construction.

Table 7: A label-free check, run on structures alone before any compound is assayed, against the coverage each design later achieved. The domain classifier separates audit compounds from batch compounds, which is the pair the predicted counts require; one half means they cannot be told apart. The block shares are the fraction of batch molecules whose building block at each position also occurs in the audit set.

	domain	block share at position			
split design	classifier	1	2	3	cov.
random assignment	0.498	1.00	1.00	1.00	98%
withhold pos. 3	0.985	1.00	1.00	0.00	48%
withhold pos. 2	0.989	1.00	0.00	1.00	37%
withhold pos. 1	0.997	0.00	1.00	1.00	35%

3.11 Calibrating against a cheaper label

The audit sections asked which compounds a round must measure. The same assumption can fail along the other axis, in what is measured on them: the quantity a campaign cares about is an affinity, and what it can afford at scale is a proxy, such as a selection count or a single-shot readout. Thresholds are then set against the proxy while the statement the team wants concerns the affinity. The compounds may be a single random sample throughout, and they are in what follows, so nothing about the population is in question; it is the labelled pair that is no longer exchangeable, because the calibration set carries the proxy's label where the audit and the batch carry the campaign's. What that costs is not a question about how often the two labels agree.

One case is free, and for the reason that makes the construction rank-based. No part of the rule reads the numerical value of a score: a threshold is an order statistic of the calibration scores of one class, a prediction set follows from comparisons against it, and every count in the outcome table is a function of ranks. So the table is unchanged by any strictly increasing transformation of the score, which may therefore be a potency, a logit or an arbitrary monotone rescaling rather than a probability; we verify this on the library screen, where three such scales give tables agreeing cell for cell (ESI Section S15). The same argument applies to the label. If the proxy is an increasing function of the affinity, then being above a proxy cut-off and being above the corresponding affinity cut-off are the same event, so the two labels coincide and nothing in the analysis moves.

Real proxies are not increasing functions of anything. What governs the shortfall is local, and follows from the construction: the threshold sits at a fixed quantile of the score distribution conditioned on the *proxy's* label, whereas the coverage achieved is the score distribution conditioned on the *true* label evaluated at that same threshold. The coverage achieved therefore decomposes into three terms,

$$\text{achieved} \;=\; (1-\alpha) \;+\; \big(\ell-(1-\alpha)\big) \;+\; \Delta, \qquad (1)$$

in which $\ell = F(\tau \mid \text{proxy out-class})$ is the realized level of the order statistic the threshold sits at and $\Delta = F(\tau \mid \text{true out-class}) - \ell$ is what the proxy costs. The middle term is the overshoot the batch correction introduces on purpose, and it is not small: the control below, whose proxy is the truth, measures $\ell = 0.924$ against a request of 0.90. Only the third term involves the proxy, and no part of it counts how many compounds the two labels disagree about. Measuring

Δ against the control rather than against the nominal level is what the agreement quoted below tests; the decomposition itself is an accounting identity. Disagreement costs coverage only insofar as it is structured with respect to the score.[1]

Three carbonic anhydrase isoforms let this be separated, because selection counts and measured inhibition constants are published for the same targets[59,60] (Table 8). No molecule in that release carries both, so the two tables cannot be joined and no real selection count can be tested against an affinity here; the proxy is instead constructed from the affinities, which is what allows the structure of its errors to be set deliberately while the compounds and the score are held fixed. What follows is therefore a demonstration of the mechanism rather than a measurement of any particular assay's fitness as a proxy. Varying only the calibration label, we compare a strictly increasing relabelling of the affinity at matched prevalence, which is the control, against two proxies that relabel the same fraction of each class: one choosing those compounds uniformly at random, the other choosing the ones the score ranks most extremely against their class. The two differ from the truth on identically many compounds in each class, so their error rates, their ranking of the truth and the prevalence they imply are all equal; they differ only in where their mistakes sit in the ranking.

They behave nothing alike. At a fifth of each class relabelled, the random proxy left class-conditional coverage at 0.933 and 0.938 against a request of 0.90, falling short in 18 of 240 class-repeats, while the score-structured proxy gave 0.721 and 0.721 and fell short in 240 of 240, at every isoform and both cut-offs. Equation 1 accounts for this, and measuring both of its sides in the same run agrees to within 0.004 on average and 0.011 at worst across all 36 conditions: that is a consistency check on the accounting rather than a prediction, and its value is that it locates the quantity, not that it is accurate. The audit-set intervals, meanwhile, covered 95.2% of consequence lines over every condition, against their nominal 95%, and are within Monte Carlo error across the table. They ask only that the audit and the batch be exchangeable, not that the rule be well calibrated, so they price whatever rule they are handed; a campaign that calibrates on a proxy keeps its predicted counts and loses its coverage guarantee, and only the audit tells it so.

That the requirement is local can be made exact, and in terms a campaign already uses. A team picks a quantile above which it will make compounds, and the rule here is of that shape, with one cut per class: at a request of $1-\alpha$ the out-class cut sits near the $1-\alpha$ quantile of the out-class scores and the in-class cut near the α quantile of the in-class scores, measured at 0.93 and 0.07 in these runs. Holding the relabelled share at a tenth of each class, so that the confusion matrix is the same in every condition, and moving only where in the within-class ranking those compounds are taken from, the damage follows the cuts and separates by class. Mistakes placed at the out-class cut displaced the out-class distributions by 0.036 and the in-class ones by 0.007; placed at the in-class cut the figures exchange, 0.007 and 0.036; placed anywhere between the two they left both at 0.007 to 0.009, which is the sampling floor. Coverage shortfalls track this, 54 of 120 class-repeats at a class's own cut against 6 to 19 elsewhere, and the audit intervals again covered 96.4% throughout. A proxy must therefore order compounds correctly in a narrow window around the quantile the campaign has chosen, and is free to misorder them everywhere else: the requirement attaches to a number the team states, not one it has to discover.

[1] Which is why a global summary of agreement, or a correlation between the proxy and the affinity, is the wrong thing to require of a proxy, and why comonotonicity is the right thing but is needed only locally: what Eq. 1 asks is that an isotonic regression[56,57] of the affinity on the proxy be increasing over the narrow score region where the thresholds fall, which the rule itself identifies as the region in which it declines to commit to a single label. Away from that region a proxy may misorder compounds freely. Fitting such a mapping needs compounds measured both ways, which a campaign can buy on a small sample; inference that corrects a cheap proxy using exactly such a small gold-standard sample is the subject of prediction-powered inference,[58] which repairs the bias where the account here instead characterizes when there is none to repair.

Table 8: Calibrating against a proxy label on three carbonic anhydrase isoforms, at affinity cut-offs of 7 and 8, 20 repeats each. The two proxies relabel the same fraction of each class and so agree with the truth equally well by any label-only summary; they differ in whether their mistakes are placed at random or where the score ranks compounds most extremely; the rows below name that placement. Coverage is class-conditional against a request of 0.90, pooled over isoforms and cut-offs; shortfall counts are given in the text rather than as a column.

	relabelled	agree	coverage		audit
proxy label	share	w/ truth	out	in	intervals
increasing in affinity	0.00	1.000	0.924	0.923	96%
at random	0.05	0.950	0.929	0.933	95%
by score	0.05	0.950	0.873	0.872	95%
at random	0.10	0.900	0.932	0.936	96%
by score	0.10	0.900	0.822	0.824	94%
at random	0.20	0.800	0.933	0.938	95%
by score	0.20	0.800	0.721	0.721	95%

3.12 What the audit must sample

The audit describes a future batch drawn from the same library as the audit itself. In simulation, when compounds share scaffold-level shifts of their predicted probability, the intervals held when the audit sampled the scaffolds being screened: coverage was 95.4–97.9% (Fig. 5, left). Auditing other scaffolds does not substitute. When the future batch came from new scaffolds with a scaffold effect of one standard deviation on the logit scale, coverage fell to 71% for false leads, 83% for missed in-class compounds and 88% for deferrals; with 1.5 standard deviations it fell to 62–75%.

Tox21 scaffolds. On Tox21 scaffold splits, calibration on other scaffolds broke the coverage guarantee: per-class coverage fell short in 28% of splits for out-class compounds and 18% for in-class compounds, two to three times the exact-law expectation (22.8 and 18.2 of 240 splits). An audit set from the screened scaffolds still reported the consequences correctly (Table 3). Calibrating instead on a random pilot sample of the screened scaffolds restored the coverage guarantee: per-class violations then followed the exact law (ESI Table S9).

Solubility. We also calibrated on all calibration compounds (about 1390) and then screened lipophilic compounds. Coverage of the soluble class among lipophilic audit compounds averaged 72% and fell below the requested 90% in all ten splits. Calibrating on the lipophilic compounds only gave

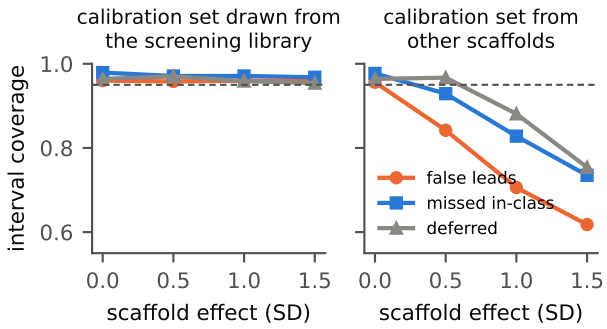


Figure 5: Coverage of 95% audit intervals when compounds are clustered in 60 scaffolds that share a shift of the predicted probability (3000 replicates per point). Left: calibration, audit and future compounds drawn from the same scaffolds, as with samples from the screening library. Right: future compounds from new scaffolds.

94% (below 90% in one split of ten, as the request allows), at the price of more measurements: 462 rather than 209 per 1000.

How large a pilot must be. In the Tox21 pilots, violations for the in-class compounds were nonetheless frequent, 29%, 13% and 12% for pilots of 200, 500 and 1000 compounds, against 26%, 12% and 9% expected, because many pilots held too few in-class compounds for the request to be attainable. At $\alpha = \delta = 0.10$ at least 22–23 calibration compounds of a class are needed, whatever the batch size; 70%, 41% and 11% of the pilots had fewer (median 13, 29 and 54 in-class compounds). A pilot of $22/\pi$ random compounds at prevalence π has an expected in-class count of exactly 22, so it leaves the request unattainable in about half of draws; the pilots run here were infeasible in 70%, 41% and 11% of draws at 200, 500 and 1000 compounds. Sizing for a stated chance of feasibility rather than for the expectation, about $28/\pi$ reaches 22 in-class compounds nine times in ten, which is 980 for NR-PPAR-γ against the 770 the expectation asks for. At that minimum the request is met only by the most conservative threshold; each further labelled compound then buys single-label calls.

Using the calibration set as its own audit. When labels are too scarce for a separate audit, the calibration set can partly serve as its own, and intervals built that way covered the future count of every cell in 97.7–100% of simulation replicates, 98.7% of realized counts on Tox21 random splits and 99.7% on the electrolyte screen (ESI Section S4). They are not a substitute, for two reasons the paper has already given. At a fixed budget of 1000 labelled compounds they are wider in every cell than splitting that budget between calibration and an audit — 21–28% for false leads and missed in-class compounds, 31–39% for confident hits and correct rejections and about twice as wide for deferrals (ESI Table S2) — while the larger calibration set buys almost nothing in return, moving the quantile the thresholds use only from 0.938 to 0.936 for a batch of 96. And being the calibration set, they inherit whatever mismatch it carries: calibrating on other scaffolds cut their coverage to 74% and on other salts to 33%. The audit set is the only one of the two that can be drawn from the chemistry the batch will contain.

4 Discussion

The audit turns a conformal classifier into a forecast: before a batch is screened, it predicts how many hits, false leads, misses and measurements the batch will contain, with intervals that hold at a stated probability. The two class-wise coverage requests move a conformal classifier through the decisions available to a screening team, and for every setting the consequences are counts in one outcome table: compounds advanced, discarded and deferred, and how many of those calls are wrong. Linked to these counts by the audit, coverage becomes a coordinate a team can steer by. An objective, usually with guarantees on what a campaign will consume and deliver that hold with a stated probability, decides where to steer. The coverage plane is where such scenarios are explored, and an audit set from the library to be screened supplies the consequences, with finite-sample valid intervals that remain valid after the team has chosen a setting by looking at them. On Tox21, AqSolDB and CALiSol-23 the intervals held at their nominal level, also when the calibration came from other scaffolds or other salts, the audit having come from the screened compounds. Each of those checks divides one labelled pool at random, however, so they verify the construction rather than establishing that it survives a separately assembled batch, and the two divisions that do test that — by provenance and by building block — break it.

Using the method in a screen. The results suggest the following workflow:

1. Label a random sample of the library to be screened and use it as the calibration set. It must contain at least 22–23 compounds of the rarer class for a request at $\alpha = \delta = 0.10$ to be attainable, and more for a useful one: keep labelling randomly chosen compounds: $22/\pi$ at prevalence π makes the request attainable only about half the time, and about $28/\pi$ reaches it nine times in ten. Choosing compounds by their score to find the rare class faster breaks exchangeability unless the selection is accounted for.[61]

2. Label a second random sample as the audit set. Its size sets the width of the intervals: for an outcome with frequency p in a batch of m, n audit compounds give a 95% interval of roughly $\pm 2m\sqrt{p(1-p)(1/n+1/m)}$, about ± 35 per 1000 for $p = 0.1$ and $n = 400$ (ESI Section S3). Choose n from the width the decision needs, not from validity, which any audit size delivers. Read the floor first, since no audit narrows an interval below $\pm 2\sqrt{mp(1-p)}$: accepting r times that floor needs $n = m/(r^2-1)$, so an audit the size of the batch is the natural place to stop and an audit much larger than the batch buys almost nothing. Batch size decides what this costs. For a 96-well plate a few hundred audit compounds already sit within a fifth of the floor, and buying more is close to wasted; for a batch of 1000 the same audit sits more than twice the floor, and closing that gap means labelling on the scale of the batch itself. With class probabilities and few labels, the calibration set can serve instead, at a cost in width that Section 3.12 quantifies at the same labelling budget.

3. Draw the audit from the chemistry the next batch will contain. Where a batch introduces building blocks,

scaffolds or series the audit does not hold, matching its composition is not enough: an audit carrying the same proportion of new chemistry drawn from different blocks performed no better than one ignoring the shift, while an audit sampling the batch's own blocks recovered the nominal level (Section 3.8). If part of the next batch is novel in this sense, a pilot of that chemistry is the indicated step, and on the electrolyte compilation it restores most of what a provenance boundary costs even when the pilot is divided from the batch non-randomly.

4. Before committing the batch, check that the calibration set, the audit set and the batch cannot be told apart. Training a classifier to separate calibration from batch compounds on structures alone needs no measurements and can be run on the proposed batch: an area under the curve near one half is consistent with exchangeability, and anything well above it means the intervals do not apply (Section 3.10). Comparing how many of the batch's scaffolds, series or building blocks occur in the audit says which part of the chemistry is responsible.

5. Sweep the coverage requests and audit each setting: intervals for missed in-class compounds, false leads and deferrals in the planned batch, with a Bonferroni correction over the sweep.

6. State a scenario, an objective and the limits that must hold with probability at least $1-\varepsilon$ (assay capacity, budget, minimum yield), and choose the best setting among those whose audit bounds satisfy the limits. If the objective and limits are monotone in the outcomes, the table of Pareto settings is all the team needs to see.

7. Screen, and compare the realized counts with the predicted intervals as a check on exchangeability.

The audit decides which questions can be asked. Calibration and audit sets are usually treated as machinery, chosen for size and then forgotten. They are better understood as the statement of what the guarantees are about, because a sample licenses claims about the population it was drawn from and about nothing else. How the audit is built therefore fixes, before any model is fitted, which questions a campaign can answer, and three choices in building it each foreclose something. The population fixes the domain: an audit drawn from the top of a ranking prices further batches drawn from the top of that same ranking, and says nothing about compounds further down, which is where a campaign goes once the top is exhausted. The unit fixes what a count means: an audit sampled one compound per series estimates the rate at which series are real and prices a batch of new series, not a batch of analogues of a few confirmed ones, whose effective size is the number of series it contains. The chemistry fixes the reach: an audit holding only building blocks already screened prices batches built from those blocks, and carrying the same proportion of new chemistry as the batch without sampling that chemistry is not a substitute (Section 3.8). The label fixes the claim: what must be exchangeable is each candidate together with its recorded label, so a campaign that sets its thresholds against a cheap readout has made a statement about that readout and not about the quantity it cares about, however interchangeable the compounds are (Section 3.11). The score belongs to the same declaration, since thresholds and the sets they define move when the model is refitted, so an audit describes the deployment it was measured under.

Reading the four together makes the condition less forbidding than it first appears, because exchangeability is sufficient but more than is needed. A coverage request concerns one quantile, and what the request actually buys is an agreement between calibration and batch at the cut it induces; a discrepancy elsewhere in the ranking is free. That is why a proxy whose errors fall away from the cut costs nothing while one whose errors fall on it is ruinous at the same error rate, and it suggests that the familiar violations might be graded the same way, by where in the ranking a population differs rather than by whether it differs at all. We have tested this only on the label, where the structure of the discrepancy can be set deliberately; doing the same for a change of chemistry would ask how much of the loss under a block holdout is attributable to the neighbourhood of the thresholds, which the designs here do not separate.

Two consequences follow for practice. The question should be settled before the audit is built, because an audit assembled for one question usually cannot be repurposed for another; a campaign that discovers halfway through that it needs to speak about a different population has to measure again. And because an audit is also an option on the questions a campaign may ask later, it is worth spending part of it on compounds a project would not otherwise measure, lower-ranked compounds or a few members of each of several series, so that when the obvious batch is exhausted the next question can still be answered. Diversity in an audit is not a statistical nicety but the purchase of future questions.

Campaigns run batch by batch. Automated and self-driving laboratories synthesize and characterize in parallel, so their campaigns are sequences of batches, each with a fixed capacity, and the framework applies to each batch as it stands. Three features of that setting make counts the right object and make an audit cheap. The constraints are physical and integer: a plate has 96 or 384 wells, a run has a reagent budget, and a round must return enough confirmed hits to justify the next. None of those is a constraint on a rate, and a scheduler cannot act on one — it needs the probability that this plate exceeds capacity, which is what a count interval states and what a risk bound does not. The batches are small and frequent, which is the regime in which the draw of the batch dominates the uncertainty, so the gap between a rate and a count is widest exactly here. And the audit is affordable because the unit of cost is the round: reserving a slice of a committed plate spends candidate slots, not an extra round. Batch campaigns also tend to pursue different goals from a search for a single optimum, though batch Bayesian optimization has long addressed the parallel setting directly. Because a batch is evaluated in parallel, its members are chosen to be diverse,[62–64] and the campaign often seeks coverage of a chemical or physical design space: mapping a landscape, as in space-filling designs and closed-loop phase mapping,[65,66] contracting onto promising regions,[67] or finding a set of optima such as a Pareto front.[68] In each of these modes a batch is a selection from a scored library under a fixed capacity, and its consequences are what the audit predicts. The batch size m is the plate or run size, and the sce-

nario constraints are the capacity of the platform, the reagent or instrument budget per run and the minimum yield needed to proceed. What the outcome table accounts for is which prediction set a candidate receives, not how the batch it belongs to was assembled. A planner that fills a plate with the highest-ranked candidates does not draw a random sample of the library, so an audit reserved from the library does not price that plate — which Section 3.9 measures rather than assumes, and which fails for a reason a team can see without any model: the plate is several times richer than the collection it was chosen from. The repair is to sample the plate and not the library, and it is exact, because a random slice of a committed batch is exchangeable with the rest of that batch however the batch was chosen. A campaign that adapts therefore reserves a random slice of every committed round, and the guarantee attaches to the remainder of that round. Methods for data chosen by the model are the relevant prior art here rather than a complementary aside.[61,69,70] When the model is retrained between batches, the thresholds change, so the calibration, the audit and the map are redone for the new deployment. Two practices follow. First, reserving a small random audit slice in every batch, measured whatever the rule decided, checks continually that the screened library still matches the calibration data; it is the acceptance sample of quality control applied to each lot. Second, the guarantee holds per batch with probability $1-\varepsilon$; a statement over all K batches of a campaign needs level ε/K for each. Because the audit bounds are stated in counts, they can be checked automatically, and the Pareto table is a compact menu from which a planner, including an AI agent, can choose a setting without a cost ratio being stated; we do not test such a planner here. What the campaign experiment adds is that the rule such a planner must follow is short and does not depend on its acquisition function: reserve a random slice of each committed round, price the remainder with it, and draw a fresh slice next round. The first clause holds because a random division of a committed batch is exchangeable with its remainder however the batch was chosen, so the planner's selection drops out of the argument entirely and no correction for it is needed; the third because one round at a sub-percent hit rate removes most of what was there to find, so an audit prices the round it came from and not its successor. A platform that follows those two clauses has finite-sample bounds on every quantity its scheduler is constrained by, for the plate it is about to commit.

Why parallel changes the economics. Automated campaigns are often pictured as sequential loops that adapt after every experiment, and a loop of that kind corrects its own errors: it proposes, measures, and the surprise moves the model, whether or not the uncertainty it reported was calibrated. Laboratories do not work that way, because the cost of an experiment is dominated by the round rather than by the candidate. A plate, a synthesis run or an instrument cycle costs nearly the same whether it is full or half empty, so cost rises in steps at each new round and barely at all within one. Three consequences follow, and they run against the intuition that smaller, more frequent batches are the way to keep a campaign honest.

The first is that an audit is cheap. When the round is the unit of cost, reserving part of a committed batch for a random sample of the library costs a few candidate slots rather than an extra round, so the measurements that price the next batch are close to free in money and are what the guarantees are bought with. The second is that correction is bought by composition rather than by cadence. Running more rounds to adapt sooner pays the fixed cost repeatedly, whereas reserving slots and spending the confidence level across a declared number of looks costs a fraction of one batch (Section 3.3). The third runs opposite to the sequential picture: the larger the batch that must be committed before anything is learned, the less opportunity there is to adapt within it, and the more of the decision rests on what was predicted beforehand. Parallelism and predicted consequences are therefore complements rather than alternatives. A laboratory that could run one cheap experiment at a time would have little use for an interval on a batch; a laboratory that must commit several hundred wells at once has the most use for one. On that reading a committed batch is an allocation under fixed capacity, divided between the candidates a setting advances, the random sample that prices the following round, and, where the next round will reach into chemistry the audit does not hold, a pilot of that chemistry (Section 3.8).

Limitations: what breaks exchangeability. All guarantees rest on exchangeability between the calibration, audit and screened candidates, each taken together with its recorded label. A change of chemical series, scaffold or assay breaks it, as the scaffold results show, and so does calibrating against a label other than the one the claim concerns, which breaks it without touching the composition of any set. So does a batch selected with the model, as in active learning, unless the selection is part of the design, and so does a library proposed by a generative model or extended into new chemistry, as the new-salt screen shows;[61,70] methods for shifted data are complementary.

Correlated candidates. Compounds that share scaffolds are correlated, so the effective size of a pilot or audit sample is closer to the number of series sampled than to the number of compounds, as the account of what an audit decides sets out above. We can now put numbers to it. Decomposing the variance of the label within and between clusters gives an intracluster correlation of 0.60 for the electrolytes at the level of the salt–solvent system, over the 138 systems that carry one, 10.8 formulations each, hence a design effect of 6.9: the 1495 room-temperature formulations are worth about 220 independent ones. For the library at the level of the first-position block the correlation is 0.35 with 680 molecules per block, a design effect of 238, so the 200 000 molecules analysed are worth about 840. What this does and does not touch is worth separating. It does not touch the audit intervals: the hypergeometric construction conditions on the total count in the audit and the batch together, which absorbs the between-cluster variance, so under random assignment both the coverage and the width of an interval are invariant to the intracluster correlation. Nor does it touch the audit-size rule, since both variance terms inflate together and the ratio r cancels the design effect. Where it bites is any statement that treats a count of candidates as a count of independent observations: the reach of a pilot, the number of clusters a sweep has actually explored, and the confidence one may place in a coverage figure estimated from correlated splits. The attainability rule for a pilot is a simple-random-

sample binomial result; a pilot of new chemistry is drawn by whole blocks, so its in-class count is overdispersed, and a sample-size rule inflates by the design effect itself rather than by its square root.[54,55] The pools used here are diverse enough for this to be immaterial, with 328 similarity clusters among 384 audited solubility compounds and over 1200 among the Tox21 compounds. It governs libraries built combinatorially or around a single target, and the DNA-encoded library is one: its random-assignment arm, which we report as the design under which the intervals hold, draws calibration, audit and batch from molecules that differ by single building blocks, so its effective sample size is smaller than its molecule count and the agreement we find there should be read as a check on the construction rather than as a measure of how much such a sample is worth.

A structural check does not substitute for a pilot. We found no relation between how well an interval covered and how similar the held-out building blocks were to those the audit held; blocks indistinguishable by fingerprint gave the lowest coverage of all, being stereoisomers whose in-class rates differ severalfold, so similarity measured on an achiral representation cannot certify a batch that differs stereochemically from its audit. The same limit bounds the pre-flight test, which inherits whatever the representation cannot see: it separated every block holdout here and the two preparations that fail, but a shift carried by features the fingerprints omit would pass it, and the provenance division of Section 3.7 is exactly such a shift. Its verdicts are therefore asymmetric, refuting an arrangement rather than certifying one, and a richer representation would widen what it can refute without ever making a pass conclusive.

The price of simultaneity, and the reach of the sweep. Each interval covers one outcome of one setting; statements about many outcomes or settings need the Bonferroni correction, which widens intervals; with 144 requests and four bounded outcomes, the joint guarantee cost about 110 measurements per 1000 at an audit of 384 compounds. Larger audits lower this price, and bounds that exploit the monotone structure of the sweep could lower it further; an envelope valid simultaneously over all thresholds is the natural instrument, and one now exists for the false discovery proportion of a conformal selection.[28] The sweep reaches every two-threshold rule only with class probabilities, and the cost-optimal rule is a two-threshold rule only if the probability of being in-class increases with the score. Cost ratios are rarely known precisely, so we report the range of ratios for which a setting is cheapest.

The datasets and their labels. Tox21 labels are missing for 7.6–26.0% of the 7592 compounds depending on the endpoint, and the missingness is not random: the 2926 compounds with no missing endpoint are active in 2.7% of their labelled cells, the 3188 missing one to three in 8.9%, and the 1478 missing four or more in 23.2%. The blanks are the inconclusive calls of the original screen, which fall disproportionately on cytotoxic and interfering compounds — the kind a real primary screen contains and a model scores high. Every Tox21 audit set here is therefore a random sample of the *labelled* subpopulation rather than of the library, at a lower in-class rate than the library carries. That is an instance of this paper's own thesis rather than an exception to it, and it means the Tox21 results should be read as verification of the construction on a clean pool rather than as evidence about a screen of the whole collection. AqSolDB combines measurements from many sources, with a typical spread of 0.5–0.7 log units between sources;[51] labels near the cut-off at $\log S = -4$ are therefore noisy, and at that spread the expected label-flip rate at the cut is 8–9% among the lipophilic compounds screened, which is the size of the coverage requests being swept; Section 3.4 measures what that costs, and the scenario is the part it costs the most; the curation rule also removed the most discordant duplicates, which a real batch would contain. CALiSol-23 likewise pools 27 sources, and several formulations were measured at many temperatures or concentrations, so neighbouring formulations are correlated. Its label also carries the provenance of the measurement rather than only the chemistry: 90% of room-temperature formulations have a single measurement in the 293–303 K window, and the temperature of that measurement is bimodal, 365 below 296 K and 917 above 300 K against 200 near 298 K, while conductivity rises by a median $0.136\,\mathrm{mS\,cm^{-1}\,K^{-1}}$, so the window spans 17% of the $8\,\mathrm{mS\,cm^{-1}}$ cut and a quarter of the formulations sit within 20% of it. Label temperature is moreover confounded with salt and with source, so the salt holdouts change the measurand as well as the chemistry. Two entries of the salt column, LiTFSI and $\mathrm{LiN(CF_3SO_2)_2}$, are the same compound under two names, 27 and 9 room-temperature rows whose mean label temperatures differ by five degrees; they were encoded separately and can therefore be split across a salt holdout. Both should be repaired before the electrolyte results are relied on: the first by interpolating each formulation to a common temperature, the second by merging the names. The guarantees refer to the labels as recorded.

Scope. We treat binary endpoints only, and one model and representation throughout. The pre-flight test, the stereoisomer blindness and the behaviour of the ranking metric at subpercent prevalence are all properties of that one fingerprint, so the design instruction of Section 3.10 is validated on a representation this section shows to be blind to a fifth of the first-position block catalogue; whether that instruction is advice or an artefact of the representation is not settled here.

5 Conclusions

The coverage level of a conformal classifier is a coordinate for navigating screening decisions, and the audited consequences, counted in compounds and bounded with a stated probability, are what a team chooses by under its own objective and guarantees. Each setting sorts a future batch into one table of prediction set by true label, and every screening quantity is a sum of its cells. A labelled audit set from the library to be screened predicts all of these counts for the next batch, with intervals that hold with a specified probability in finite samples, under exchangeability of the audit with the batch alone and after a setting is chosen. The same request led to very different consequences across twelve Tox21 endpoints, and one model offered more than a hundred distinct trade-offs for a solubility screen and for an electrolyte screen. When the calibration did not match the screened pop-

ulation, as for a new lithium salt, the audit still bounded the consequences correctly — though that check, like every positive coverage result here, divides one labelled pool at random, and the two divisions that test the intervals against a genuinely separate batch, by provenance and by building block, break them. In automated campaigns, which run in parallel batches, the same construction applies to each batch, and the bounds become capacity constraints a planner can check, provided the slice it reserves comes from the committed round and not from the library: a plate chosen from the top of a ranking was 38 times richer than the pool it came from, and an audit of that pool held nothing in-class to measure. One round also depletes what remains, so the slice is re-drawn each round rather than amortized across the campaign. Inverted, the same intervals size a round: the batch an audited count implies for a target number of hits is larger than dividing the target by the rate suggests, by a factor that grows as the rate falls, and the audit that pins the rate down is the measurement with the largest effect on what the round costs. What the audit must contain is decided by the batch it prices, since an audit that reproduces a batch's composition without sampling its chemistry prices nothing; buying a pilot of the screened population does repair it, and does so under divisions that are not random and could have failed. With an audit built this way, a team sees those consequences before the campaign starts.

Author contributions

P.H.Z. conceived the study, developed the methodology, performed the analysis and wrote the manuscript.

Conflicts of interest

There are no conflicts to declare.

Data availability

Tox21 is available through MoleculeNet, AqSolDB through its dataset reference[51] and CALiSol-23 from the DTU data repository (DOI 10.11583/DTU.c.6929599).[52] The carbonic anhydrase selection counts and affinities are from the CA-DEL release[59] (Zenodo record 17656024, DOI 10.5281/zenodo.17656024, CC BY 4.0), whose affinity values derive from ChEMBL[60] and carry CC BY-SA 3.0. The DNA-encoded library is BELKA,[53] available from Leash Biosciences through the Polaris record `leash-bio/belka-v1` under CC BY-NC-SA 4.0; the million-molecule sample used here is reproduced by the seeded digest rule given in ESI Section S14. The scripts that reproduce every table and figure are provided with the ESI, together with the processed data, split indices and stored model scores. They implement the finite-sample threshold corrections, randomized tie-breaking, outcome tables, prediction intervals for future batches from an audit set or from the calibration set alone, coverage-request sweeps with corrections for selection, and the cost-ratio analysis. They are released under the MIT licence at `https://github.com/phzwart/conformal-screening`. [Tagged version and archived Zenodo DOI, on acceptance.]

Acknowledgements

This work was supported by the U.S. Department of Energy, Office of Science, under Contract No. DE-AC02-05CH11231 as part of the Laboratory Directed Research and Development (LDRD) program. Additional support came in part from the U.S. Department of Energy, Office of Science, Scientific Discovery through Advanced Computing (SciDAC) program FORUM-AI. Generative AI was used in preparing this manuscript: for drafting and editing prose, for writing and refactoring analysis code, and for assembling tables and figures from computed results. The author designed the study, chose the analyses, and has verified the reported numbers against the stored outputs of the scripts that produced them.

References

[1] Vladimir Vovk, Alex Gammerman, and Glenn Shafer. *Algorithmic Learning in a Random World*. Springer, New York, 2005. doi: 10.1007/b106715.

[2] Anastasios N. Angelopoulos and Stephen Bates. Conformal prediction: A gentle introduction. *Foundations and Trends in Machine Learning*, 16(4):494–591, 2023. doi: 10.1561/2200000101.

[3] Stephen Bates, Anastasios Angelopoulos, Lihua Lei, Jitendra Malik, and Michael I. Jordan. Distribution-free, risk-controlling prediction sets. *Journal of the ACM*, 68(6), September 2021. ISSN 0004-5411. doi: 10.1145/3478535. URL `https://doi.org/10.1145/3478535`.

[4] Anastasios N. Angelopoulos, Stephen Bates, Adam Fisch, Lihua Lei, and Tal Schuster. Conformal risk control. In *International Conference on Learning Representations*, 2024.

[5] Anastasios N. Angelopoulos, Stephen Bates, Emmanuel J. Candès, Michael I. Jordan, and Lihua Lei. Learn then test: calibrating predictive algorithms to achieve risk control. *arXiv preprint arXiv:2110.01052*, 2021.

[6] Lewis H. Mervin, Simon Johansson, Elizaveta Semenova, Kathryn A. Giblin, and Ola Engkvist. Uncertainty quantification in drug design. *Drug Discovery Today*, 26(2):474–489, 2021. doi: 10.1016/j.drudis.2020.11.027.

[7] A. R. Thatcher. Relationships between Bayesian and confidence limits for predictions. *Journal of the Royal Statistical Society: Series B (Methodological)*, 26(2):176–192, 1964. doi: 10.1111/j.2517-6161.1964.tb00551.x.

[8] William Q. Meeker, Gerald J. Hahn, and Luis A. Escobar. *Statistical Intervals: A Guide for Practitioners and Researchers*. Wiley, Hoboken, NJ, 2nd edition, 2017.

[9] David Lopez-Paz and Maxime Oquab. Revisiting classifier two-sample tests. In *International Conference on Learning Representations*, 2017.

[10] Ulf Norinder, Lars Carlsson, Scott Boyer, and Martin Eklund. Introducing conformal prediction in predictive modeling. A transparent and flexible alternative to applicability domain determination. *Journal of Chemical Information and Modeling*, 54(6):1596–1603, 2014.

[11] Ulf Norinder and Scott Boyer. Binary classification of imbalanced datasets using conformal prediction. *Journal of Molecular Graphics and Modelling*, 72:256–265, 2017.

[12] Jiangming Sun, Lars Carlsson, Ernst Ahlberg, Ulf Norinder, Ola Engkvist, and Hongming Chen. Applying Mondrian cross-conformal prediction to estimate prediction confidence on large imbalanced bioactivity data sets. *Journal of Chemical Information and Modeling*, 57(7):1591–1598, 2017.

[13] Fredrik Svensson, Ulf Norinder, and Andreas Bender. Modelling compound cytotoxicity using conformal prediction and PubChem HTS data. *Toxicology Research*, 6(1):73–80, 2017.

[14] Nicolas Bosc, Francis Atkinson, Eloy Felix, Anna Gaulton, Anne Hersey, and Andrew R. Leach. Large scale comparison of QSAR and conformal prediction methods and their applications in drug discovery. *Journal of Cheminformatics*, 11:4, 2019.

[15] Isidro Cortés-Ciriano and Andreas Bender. Concepts and applications of conformal prediction in computational drug discovery. *arXiv preprint arXiv:1908.03569*, 2019.

[16] Andrea Morger, Fredrik Svensson, Staffan Arvidsson McShane, Niharika Gauraha, Ulf Norinder, Ola Spjuth, and Andrea Volkamer. Assessing the calibration in toxicological in vitro models with conformal prediction. *Journal of Cheminformatics*, 13:35, 2021. doi: 10.1186/s13321-021-00511-5.

[17] Gökçe Geylan, Leonardo De Maria, Ola Engkvist, Florian David, and Ulf Norinder. A methodology to correctly assess the applicability domain of cell membrane permeability predictors for cyclic peptides. *Digital Discovery*, 3:1761–1775, 2024. doi: 10.1039/D4DD00056K.

[18] Staffan Arvidsson McShane, Ulf Norinder, Jonathan Alvarsson, Ernst Ahlberg, Lars Carlsson, and Ola Spjuth. CPSign: conformal prediction for cheminformatics modeling. *Journal of Cheminformatics*, 16:75, 2024. doi: 10.1186/s13321-024-00870-9.

[19] Siddhartha Laghuvarapu, Zhen Lin, and Jimeng Sun. CoDrug: conformal drug property prediction with density estimation under covariate shift. In *Advances in Neural Information Processing Systems*, volume 36, 2023.

[20] Ruben Buendia, Thierry Kogej, Ola Engkvist, Lars Carlsson, Henrik Linusson, Ulf Johansson, Paolo Toccaceli, and Ernst Ahlberg. Accurate hit estimation for iterative screening using Venn–ABERS predictors. *J. Chem. Inf. Model.*, 59(3):1230–1237, 2019. doi: 10.1021/acs.jcim.8b00724.

[21] Ruben Buendia, Ola Engkvist, Lars Carlsson, Thierry Kogej, and Ernst Ahlberg. Venn–Abers predictors for improved compound iterative screening in drug discovery. *Proc. Mach. Learn. Res.*, 91:201–219, 2018.

[22] Fredrik Svensson, Avid M. Afzal, Ulf Norinder, and Andreas Bender. Maximizing gain in high-throughput screening using conformal prediction. *Journal of Cheminformatics*, 10:7, 2018. doi: 10.1186/s13321-018-0260-4.

[23] Vladimir Vovk. Conditional validity of inductive conformal predictors. *arXiv*, 2012. doi: 10.48550/arXiv.1209.2673. Extended version of ACML 2012 paper.

[24] Michelle Bian and Rina Foygel Barber. Training-conditional coverage for distribution-free predictive inference. *Electronic Journal of Statistics*, 17(2):2044–2066, 2023. doi: 10.1214/23-EJS2145.

[25] Paulo C. F. Marques. Universal distribution of the empirical coverage in split conformal prediction. *Statistics & Probability Letters*, 219:110350, 2025. ISSN 0167-7152. doi: 10.1016/j.spl.2024.110350. URL `https://www.sciencedirect.com/science/article/pii/S0167715224003195`.

[26] S. S. Wilks. Determination of sample sizes for setting tolerance limits. *The Annals of Mathematical Statistics*, 12(1):91–96, 1941.

[27] Sangdon Park, Osbert Bastani, Nikolai Matni, and Insup Lee. PAC confidence sets for deep neural networks via calibrated prediction. In *International Conference on Learning Representations*, 2020.

[28] Ziang Song, Ying Jin, and Emmanuel J. Candès. Everywhere valid bounds on false discovery proportions in conformal inference. *arXiv preprint arXiv:2605.20726*, 2026.

[29] Jelle J. Goeman and Aldo Solari. Multiple testing for exploratory research. *Statistical Science*, 26(4):584–597, 2011. doi: 10.1214/11-STS356.

[30] Eugene Katsevich and Aaditya Ramdas. Simultaneous high-probability bounds on the false discovery proportion in structured, regression and online settings. *The Annals of Statistics*, 48(6):3465–3487, 2020. doi: 10.1214/19-AOS1938.

[31] Gilles Blanchard, Pierre Neuvial, and Etienne Roquain. Post hoc confidence bounds on false positives using reference families. *The Annals of Statistics*, 48(3):1281–1303, 2020. doi: 10.1214/19-AOS1847.

[32] Ying Jin and Emmanuel J. Candès. Selection by prediction with conformal p-values. *Journal of Machine Learning Research*, 24(244):1–41, 2023.

[33] Mauricio Sadinle, Jing Lei, and Larry Wasserman. Least ambiguous set-valued classifiers with bounded error levels. *Journal of the American Statistical Association*, 114(525):223–234, 2019.

[34] Jordan Lekeufack, Anastasios N. Angelopoulos, Andrea Bajcsy, Michael I. Jordan, and Jitendra Malik. Conformal decision theory: Safe autonomous decisions from imperfect predictions. In *2024 IEEE International Conference on Robotics and Automation (ICRA)*, pages 11668–11675, 2024.

[35] Shayan Kiyani, George Pappas, Aaron Roth, and Hamed Hassani. Decision theoretic foundations for conformal prediction: Optimal uncertainty quantification for risk-averse agents. In *Proceedings of the 42nd International Conference on Machine Learning*, volume 267 of *Proceedings of Machine Learning Research*, 2025.

[36] C. K. Chow. On optimum recognition error and reject tradeoff. *IEEE Transactions on Information Theory*, 16(1):41–46, 1970. doi: 10.1109/TIT.1970.1054406.

[37] Foster Provost and Tom Fawcett. Robust classification for imprecise environments. *Machine Learning*, 42(3):203–231, 2001.

[38] Tadeusz Pietraszek. Optimizing abstaining classifiers using ROC analysis. In *Proceedings of the 22nd International Conference on Machine Learning*, pages 665–672, 2005.

[39] Pascal Massart. The tight constant in the dvoretzky–kiefer–wolfowitz inequality. *Annals of Probability*, 18(3):1269–1283, 1990. doi: 10.1214/aop/1176990746.

[40] A. Charnes and W. W. Cooper. Chance-constrained programming. *Management Science*, 6(1):73–79, 1959.

[41] Philip B. Stark. Conservative statistical post-election audits. *The Annals of Applied Statistics*, 2(2):550–581, 2008. doi: 10.1214/08-AOAS161.

[42] Mark Lindeman and Philip B. Stark. A gentle introduction to risk-limiting audits. *IEEE Security & Privacy*, 10(5):42–49, 2012. doi: 10.1109/MSP.2012.56.

[43] Harold F. Dodge and Harry G. Romig. *Sampling Inspection Tables: Single and Double Sampling*. Wiley, New York, 1944.

[44] Tilmann Gneiting, Fadoua Balabdaoui, and Adrian E. Raftery. Probabilistic forecasts, calibration and sharpness. *J. R. Stat. Soc. B*, 69(2):243–268, 2007. doi: 10.1111/j.1467-9868.2007.00587.x.

[45] Greg Landrum. RDKit: Open-source cheminformatics. `https://www.rdkit.org`, 2025. [Online; accessed 2026-02-13].

[46] David Rogers and Mathew Hahn. Extended-connectivity fingerprints. *Journal of Chemical Information and Modeling*, 50 (5):742–754, 2010. doi: 10.1021/ci100050t.

[47] Liudmila Prokhorenkova, Gleb Gusev, Aleksandr Vorobev, Anna Veronika Dorogush, and Andrey Gulin. CatBoost: unbiased boosting with categorical features. In *Advances in Neural Information Processing Systems*, volume 31, pages 6638–6648, 2018.

[48] Ruili Huang, Menghang Xia, Dac-Trung Nguyen, Jianying Zhao, Srilatha Sakamuru, Tingjun Zhao, Maxine Na, Swati A. Shahane, Anastasia Rossoshek, and Anton Simeonov. The tox21 data challenge to build predictive models of nuclear receptor and stress response pathways. *Nature Biotechnology*, 34(8):828–837, 2016. doi: 10.1038/nbt.3659.

[49] Andreas Mayr, Günter Klambauer, Thomas Unterthiner, and Sepp Hochreiter. Deeptox: Toxicity prediction using deep learning. *Frontiers in Environmental Science*, 3:80, 2016. doi: 10.3389/fenvs.2015.00080.

[50] Guy W. Bemis and Mark A. Murcko. The properties of known drugs. 1. Molecular frameworks. *Journal of Medicinal Chemistry*, 39(15):2887–2893, 1996. doi: 10.1021/jm9602928.

[51] Murat Cihan Sorkun, Abhishek Khetan, and Süleyman Er. Aqsoldb: a curated reference set of aqueous solubility and 2d descriptors for a diverse set of compounds. *Scientific Data*, 6: 143, 2019. doi: 10.1038/s41597-019-0151-1. URL `https://www.nature.com/articles/s41597-019-0151-1`.

[52] Paolo de Blasio, Jonas Elsborg, Tejs Vegge, Eibar Flores, and Arghya Bhowmik. CALiSol-23: Experimental electrolyte conductivity data for various Li-salts and solvent combinations. *Scientific Data*, 11(1):750, 2024.

[53] Ian Quigley, Andrew Blevins, Nate Wilkinson, Maxwell Kleinsasser, Victor Shlepov, Vlad Vinogradov, and Jonathan Hsu. BELKA: the Big Encoded Library for Chemical Assessment. In *Advances in Neural Information Processing Systems 37, Competition Track*, 2024. URL `https://openreview.net/forum?id=zwppB4butE`. Dataset: Leash Biosciences; Polaris record `leash-bio/belka-v1`, CC BY-NC-SA 4.0.

[54] William G. Cochran. *Sampling Techniques*. John Wiley & Sons, New York, 3rd edition, 1977.

[55] Leslie Kish. *Survey Sampling*. John Wiley & Sons, New York, 1965.

[56] Miriam Ayer, H. D. Brunk, G. M. Ewing, W. T. Reid, and Edward Silverman. An empirical distribution function for sampling with incomplete information. *The Annals of Mathematical Statistics*, 26(4):641–647, 1955. doi: 10.1214/aoms/1177728423.

[57] Bianca Zadrozny and Charles Elkan. Transforming classifier scores into accurate multiclass probability estimates. In *Proceedings of the Eighth ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pages 694–699, 2002. doi: 10.1145/775047.775151.

[58] Anastasios N. Angelopoulos, Stephen Bates, Clara Fannjiang, Michael I. Jordan, and Tijana Zrnic. Prediction-powered inference. *Science*, 382(6671):669–674, 2023. doi: 10.1126/science.adi6000.

[59] Yuxuan He, Duanhua Cao, Xiaozhe Tan, Mingan Gao, Xiaojun Yao, Shuangjia Gu, and Pheng-Ann Heng. CA-DEL: an open multi-target, multi-modal benchmark for learning from DNA-encoded library screens, 2026. Data: Zenodo record 17656024, *CA_DEL_database*, CC BY 4.0.

[60] Barbara Zdrazil, Eloy Felix, Fiona Hunter, Emma J. Manners, James Blackshaw, Sybilla Corbett, Marleen de Veij, Harris Ioannidis, David Mendez Lopez, Juan F. Mosquera, Maria Paula Magarinos, Nicolas Bosc, Ricardo Arcila, Tevfik Kizilören, Anne Gaulton, A. Patrícia Bento, Melissa F. Adasme, Peter Monecke, Gregory A. Landrum, and Andrew R. Leach. The ChEMBL database in 2023: a drug discovery platform spanning multiple bioactivity data types and time periods. *Nucleic Acids Research*, 52(D1):D1180–D1192, 2024. doi: 10.1093/nar/gkad1004.

[61] Rina Foygel Barber, Emmanuel J. Candès, Aaditya Ramdas, and Ryan J. Tibshirani. Conformal prediction beyond exchangeability. *The Annals of Statistics*, 51(2):816–845, 2023.

[62] Javier González, Zhenwen Dai, Philipp Hennig, and Neil Lawrence. Batch Bayesian optimization via local penalization. In *Proceedings of the 19th International Conference on Artificial Intelligence and Statistics*, volume 51 of *Proceedings of Machine Learning Research*, pages 648–657, 2016.

[63] José Miguel Hernández-Lobato, James Requeima, Edward O. Pyzer-Knapp, and Alán Aspuru-Guzik. Parallel and distributed Thompson sampling for large-scale accelerated exploration of chemical space. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1470–1479, 2017.

[64] Benjamin J. Shields, Jason Stevens, Jun Li, Marvin Parasram, Farhan Damani, Jesus I. Martinez Alvarado, Jacob M. Janey, Ryan P. Adams, and Abigail G. Doyle. Bayesian reaction optimization as a tool for chemical synthesis. *Nature*, 590:89–96, 2021. doi: 10.1038/s41586-021-03213-y.

[65] V. Roshan Joseph. Space-filling designs for computer experiments: A review. *Qual. Eng.*, 28(1):28–35, 2016. doi: 10.1080/08982112.2015.1100447.

[66] A. Gilad Kusne, Heshan Yu, Changming Wu, Huairuo Zhang, Jason Hattrick-Simpers, Brian DeCost, Suchismita Sarker, Corey Oses, Cormac Toher, Stefano Curtarolo, Albert V. Davydov, Ritesh Agarwal, Leonid A. Bendersky, Mo Li, Apurva Mehta, and Ichiro Takeuchi. On-the-fly closed-loop materials discovery via Bayesian active learning. *Nat. Commun.*, 11:5966, 2020. doi: 10.1038/s41467-020-19597-w.

[67] David Eriksson, Michael Pearce, Jacob Gardner, Ryan D. Turner, and Matthias Poloczek. Scalable global optimization via local Bayesian optimization. In *Advances in Neural Information Processing Systems*, volume 32, 2019.

[68] Samuel Daulton, Maximilian Balandat, and Eytan Bakshy. Differentiable expected hypervolume improvement for parallel multi-objective Bayesian optimization. In *Advances in Neural Information Processing Systems*, volume 33, 2020.

[69] Ryan J. Tibshirani, Rina Foygel Barber, Emmanuel J. Candès, and Aaditya Ramdas. Conformal prediction under covariate shift. In *Advances in Neural Information Processing Systems 32*, pages 2530–2540, 2019.

[70] Clara Fannjiang, Stephen Bates, Anastasios N. Angelopoulos, Jennifer Listgarten, and Michael I. Jordan. Conformal prediction under feedback covariate shift for biomolecular design. *Proceedings of the National Academy of Sciences*, 119(43): e2204569119, 2022. doi: 10.1073/pnas.2204569119. URL `https://arxiv.org/abs/2202.03613`.